

Lecture 1-2

Table of contents

Overview	1
Software	2
In Summary: Following on from Tukey's Steps	3
A word of caution on practical data analysis	3
Datasets and Variable Types	10
Dataset resources	10
Dataset types - Measurement levels	11
Dataset file formats	13
Types of EDA	13
EDA methods - Summary Statistics	14
EDA methods - Plots	14
Distribution of a Single Quantitative Variable	15
Visual Display of a Single Qualitative Variable	15
R Examples	16
Example 1 - Female headed households in SA	16
Example 2 - Gapminder	22
Class exercise	40

Overview

💡 Reading

What is EDA?

Secondary Analysis of Electronic Health Records Book Chapter 15 - Exploratory Data Analysis

Tukey's EDA book and terminology

Exploratory Data Analysis (EDA) is one of the key preliminary steps any data scientist who deals with data needs to perform before doing any analysis.

EDA is performed to understand your data, to unearth the underlying structure, to assess the quality of your data by means of summary statistics and visual representations, to discover the main attributes and characteristics of each variable in your data, to discover the relationships between your variables in your data and of course to gain insights before performing any complex modelling.

EDA is not done because it has to be done or because it is one of the requirements of a data science project, it is done for scientific reasoning, not housekeeping and to discover what you don't yet know you should model. Of course, as in any project, you should start your thinking process around the hypotheses you might already have, the questions you want to answer, however, EDA is an open-ended, iterative process. It needs to be [reproducible](#). Do not save output!

Software

In this course, we will only use [R](#) and teach in R. We assume you know R or learned R in the STA5075Z - Statistical Computing in R course. I need to mention that R is just a tool, and there are several other tools to use such as [Python](#). EDA is not about learning R, or making plots in R etc. — it's about thinking statistically with data to help modelling and decision-making, regardless of the tool.

Why R?

- Both R and Python are free.
- R already has all of the statistics support because it was developed by statisticians for statisticians. A lot of statistical modelling research is conducted in R.
- Python was originally developed as a programming language for software development, DS tools (scikit-learn, pandas, numpy) were added on. Though the majority of DL research is done in Python, such as keras, PyTorch.

- R has Tidyverse, a set of packages that makes it easy to import, manipulate, visualise and report data.
- Very easy to generate dashboards using R Shiny.
- It is the language I know the best, I know very little Python.
- Python and R programmers get inspired from each other, ie. Python's plotnine inspired by R's ggplot2, and R's rvest by Python's BeautifulSoup.
- You can also use functions written in Python with python function in R.
- You can run R code from Python with rp2 package, and you can run Python code from R using reticulate. R version of DL package Keras calls Python.
- Though, I do encourage you to learn Python as well. No harm in two languages.

Take away: There is no winner, you are here to learn the skills, your focus should be on skills. If you can program in R, you can do it in any other language.

Reading

[Python vs. R for Data Science: What's the Difference \(by Richie Cotton\)](#)
[R vs Python for Data Science: The Winner is... \(by Martijn Theuwissen\) }](#)

In Summary: Following on from Tukey's Steps

Tukey's EDA book provides techniques and advice about how to explore data.

- The approach of EDA is detective in character, it is a search for clues. Some of the clues may be misleading, but some will lead to discoveries.
 - Tukey favors simplicity because simple statements are clear.
 - Tukey favors clear visual displays of quantitative facts.
 - Tukey likes precision, it is far better to be able to say some response measure is a linear function of a particular stimulus variable than to say it increases with the stimulus variable.
 - Tukey favors depth of analysis. It is always good to look at the residuals.
 - Tukey values accuracy. A misplaced decimal vs a misplaced digit.
 - Tukey values replicability of summary observations in situations containing aberrant observations.
-

A word of caution on practical data analysis

One of the main aims of this course is to put you in a position of being able to perform the multivariate statistical analysis of your own research projects, in whatever field this may be.

Most of the examples used in these notes are themselves real-world studies, and so you will get some idea of some of the complexities involved in gathering and analysing data. Having said that, there is an obvious need in an introductory course like this one to choose data sets that work' and that can be used to illustrate the techniques. We therefore do not discuss many of the practical difficulties which inevitably arise when doing your own original research. As a result when these difficulties arise when it comes to doing your own research, you may look back on this course and think why weren't we taught that?' Unfortunately, the kinds of problems that can arise are so varied and require such different solutions that it is not possible to teach in a course such as this one. As Bartholemew et al. put it, only when one has a clear idea of where one is going is it possible to know the important questions which arise". However, the following broad areas should be borne in mind whenever conducting an original analysis.

Missing Data

Missing data can cause severe problems for many of the techniques we will consider. Most techniques will simply drop cases which possess missing data on *any* of the variables to be included in the analysis. When the number of variables is large, as is often the case in multivariate analyses, this can result in a substantial proportion of the sample being dropped. This proportion should always be noted early in the analysis. Another critical question to ask is "why is the data missing?" and "does the missing data introduce any bias into the results?" Often, it is the people with the most extreme views that turn up as missing data by refusing to answer certain questions, which is clearly biasing. Possible solutions are *mean replacement* or other *imputation* (replacement) techniques, but these are beyond the scope of this course.

Sample Sizes

It is a general rule that the bigger the model you fit, the greater the number of cases you need. In univariate analysis and simple hypothesis testing, the calculation of required' sample sizes is reasonably straightforward, but in multivariate analysis there are only very rough guidelines where any exist at all. As a *very* rough guideline, most techniques require at least 10 respondents per parameter estimated. That means that in order to estimate a regression model with four independent variable, you need at least 50 respondents (not forgetting the constant term β_0 , there are 5 parameters to be estimated). When sample sizes are small, one should be very careful about drawing strong conclusions. This is a particular problem in student research, where sample sizes are typically very small.

Transformations

Many statistical techniques assume that data are normally distributed. Although it is again beyond the scope of this course, it is often possible to transform data that is not normally distributed into something that *is* normally distributed by using some kind of transforming function. Taking the logarithm of a set of numbers, for example, often works, as does taking the square (both of these transformations work by sucking in' the tails of the non-normal distributions). Where transformations do not help, the analyst must make a decision about whether the data is approximately normal' or 'normal enough' to continue, or whether it is necessary to use other methods (like non-parametric statistics, which tend to be harder to use but do not make any distributional assumptions).

Very often it is more convenient to look at some transform of the original variable. If the distribution is far from symmetrical, one end of the distribution will be too crowded to permit careful inspection.

Tukey deals extensively with scale transformations. He gives three main reasons for transformations:

- A transformation may be selected to produce a symmetrical distribution,
 - A transformation may increase the similarity of the spread of the different sets of numbers,
 - A transformation may straighten out a line.
-

Types of transformations

The transformations discussed range on:

- x^n
- x^{n-1}
- x^2
- $\log x$
- $-\frac{1}{x}$
- $-\frac{1}{x^2}$
- $-\frac{1}{x^n}$

Other dependent variables, e.g. counts and latencies, are occasionally transformed by taking the square root, the logarithm, or the reciprocal.

Standardisation of data

When analysing numerical data, it often happens that different variables are measured on scales of very different sizes. For example, in the above matrix the first question might ask one how many children one has and the second question might ask for one's income in Rands. Clearly, the scale of possible values for the first question (between 0 and perhaps 15) is *much* smaller than for the second (between 0 and perhaps several million Rand). For reasons that will become clearer later on, this can cause enormous problems in some multivariate techniques by giving too much influence to the variables measured on larger scales. In order to put all variables on an equal footing, it is often necessary to *standardise* the data. Because several techniques require standardised data we consider it in this introductory chapter, but it is important to realise that *not all* the techniques need the data to be standardised. Moreover, in cases where all numerical variables are measured on the same scale (e.g.all on a 1 to 5 Likert rating scale) there will be no need to standardise either.

There are several different ways to standardise data. We will illustrate the standardisation of a data matrix using the following example. Suppose that information on three variables (income, number of children, and age) has been collected from five individuals. The data is contained in the following table.

Table 1: Unstandardized Data with Summary Statistics

Person	Income	No Children	Age
a	10000	0	40
b	0	3	23
c	300000	2	32
d	150000	2	35
e	1000000	1	58
\bar{x}	292000	1.6	37.6
s	414210	1.140	12.973

where we use the usual mathematical notation \bar{x} to denote the mean and s to denote the standard deviation. Note that the variables are measured on very different scales. To standardise the data, we simply follow the steps above. For example, the standardised income of person a is given by

$$\frac{10\,000 - 292\,000}{414\,210} = -0.681$$

to three decimal places. Similarly the standardised number of children for person d is given by $(2 - 1.6)/1.140 = 0.351$. You can check for yourself that the *new* column means and standard deviations are all zero and one respectively. Since the mean of all the variables is zero, it is

possible to see at a glance which observations are below average (those that are negative) and which are above average (those that are positive).

```
# 1. Create the data frame for standardized values
std_data <- data.frame(
  Person = c("a", "b", "c", "d", "e"),
  Income = c(-0.681, -0.705, 0.019, -0.343, 1.709),
  N.Children = c(-1.403, 1.228, 0.351, 0.351, -0.526),
  Age = c(0.185, -1.125, -0.432, -0.200, 1.572)
)
```

```
kableExtra::kable(std_data)
```

Table 2: Standardised Data

Person	Income	N.Children	Age
a	-0.681	-1.403	0.185
b	-0.705	1.228	-1.125
c	0.019	0.351	-0.432
d	-0.343	0.351	-0.200
e	1.709	-0.526	1.572

The relevance of standardising data may not seem clear to you at the moment. Just bear this section in mind as you continue through the notes and refer back to it when the issue of standardisation reappears.

```
X <- matrix (c(10000,0,300000,150000,1000000,0,3,2,2,1,40,23,32,35,58),
              ncol=3,
              dimnames=list(c("a","b","c","d","e"),
                            c("Income","No Children","Age")))
```

In R we can create a matrix with the `matrix()` function. The values in the matrix are concatenated with the operator `c()`. Notice that the values needs to be entered column wise by default. The names for the two dimensions are specified by `dimnames=list("row names", "column names")`. Notice below that the row names appear to the left.

They are text, but are not part of the CONTENT of the matrix. The matrix X:5 × 3 contains only numeric values.

```
X
```

	Income	No Children	Age
a	10000	0	40
b	0	3	23
c	300000	2	32
d	150000	2	35
e	1000000	1	58

To calculate the means we apply to X, column wise (indicated by 2; 1 for row wise) the function `mean()`.

```
xbar <- apply(X,2,mean)  
xbar
```

Income	No Children	Age
292000.0	1.6	37.6

Similarly, the function `sd()` is applied to each column to calculate the standard deviations.

```
s <- apply(X,2,sd)  
s
```

Income	No Children	Age
4.142101e+05	1.140175e+00	1.297305e+01

Any numeric calculations can be performed by simply typing the expression at the R command prompt “>”.

```
(10000-292000)/414210
```

```
[1] -0.6808141
```

R has the ability to operate on a whole vector (or matrix) at once. Here the standardised values for Age is calculated by subtracting the mean from the values in column 2 and dividing resulting “column minus mean” by the standard deviation.

```
(X[,2]-1.6)/1.14
```

```
a           b           c           d           e  
-1.4035088  1.2280702  0.3508772  0.3508772 -0.5263158
```

The expressions above is simply for illustration purposes. The function `scale()` performs all the standardisation calculations in a single step. The output is again a matrix of size 5×3 , but additional attributes are provided: first the mean called “`scaled:center`”, then the standard deviations called “`scaled:scale`”.

```
scale(X)
```

```
Income No Children      Age  
a -0.68081393 -1.4032928  0.1849989  
b -0.70495627  1.2278812 -1.1254101  
c  0.01931387  0.3508232 -0.4316641  
d -0.34282120  0.3508232 -0.2004155  
e  1.70927752 -0.5262348  1.5724908  
attr(,"scaled:center")  
Income No Children      Age  
292000.0    1.6        37.6  
attr(,"scaled:scale")  
Income No Children      Age  
4.142101e+05 1.140175e+00 1.297305e+01
```

Also, be forewarned that you need to know basic statistics to be able to follow this course and if you are rusty in your first year statistics knowledge, then please read STA1000 book to refreshen your background. The aim of this course and these notes is to cover some of the more popular methods for exploring multivariate data. The perspective that we will take when looking at these techniques will be to use the minimum amount of mathematics necessary for a solid understanding of the techniques and their interpretation. However, this does not mean “no mathematics”! Over the past twenty or so years, modern statistical software packages have made it possible to run all of the techniques that we’ll cover in this course with a few clicks of a mouse, without knowing a single bit of mathematics and almost nothing about how the techniques themselves work. Clicking a mouse might give you results, but it is very difficult to know whether these results are reliable unless you know something about the underlying technique and what potential pitfalls exist. All statistical techniques, and particularly the multivariate ones, make some assumptions about the type and amount of data that should be collected and the aims of the researcher. If these are ignored, the results may not just be incorrect but misleading. In this case it would be better to put the output of an analysis in a

rubbish bin than into a report or on a manager's desk. To get this understanding, a certain amount of mathematics is needed.

Having said that, the focus of the course is on the practical use and interpretation of the techniques in the analysis of real-world examples. The kind of statistics and mathematics that will be used includes the following topics that have been covered in previous courses:

- Basic descriptive statistics (means, variances, absolute and relative frequencies, correlation)
- Hypothesis testing (z -test, t -test, F -test, χ^2 -test of association)
- One-way analysis of variance and multiple linear regression
- Two-way crosstables (contingency tables), and their analysis using the χ^2 test of association
- Use of basic mathematical notation (summation' notation, vectors, matrices)

If you are unfamiliar with any of this material, it is important to go back and revise in the first few weeks of the course.

Datasets and Variable Types

In this course, we will look at published MSc Data Science theses from [OpenUCT](#) (browse by department, type “Department of Statistical Sciences” in the search) and we will use publicly available datasets from various resources.

Dataset resources

- [Zindi Competitions](#):
- Zindi is the first data science competition platform in Africa.
- Zindi hosts an entire data science ecosystem of scientists, engineers, academics, companies, NGOs, governments and institutions focused on solving Africa’s most pressing problems.
- [Kaggle competitions](#)
- [UCT Open Data Portal - ZivaHub](#)

- StatLib Datasets Archive
 - [UC Irvine Machine Learning Repository](#)
 - The Humanitarian Data Exchange
 - StatSci.org datasets
 - City of Cape Town open data portal
 - JSE data archive
 - Vanderbilt Biostatistics Datasets
 - PhysioBank databases
 - R Data Sources for Regression Analysis
 - FiveThirtyEight data
 - TidyTuesday
 - World Health Organization
 - The National Bureau of Economic Research
 - International Monetary Fund
 - General Social Survey
 - United Nations Data
 - United Nations Statistics Division
 - European Statistics
 - UNICEF
 - CDC
 - World Bank
-

Dataset types - Measurement levels

At this point it is probably worth spending a little time discussing different data types. There are two main types of data that we need to distinguish between: *numerical* variables, and *categorical* variables.

! Important

Numerical variables

Numerical variables are measurements that can be recorded on a quantitative scale where the intervals between two values on the scale have some meaning. Essentially, this means that (a) the variable contains numbers rather than words or symbols, (b) the gaps between two numbers have some actual meaning. Examples of numerical variables are height, age, and number of children.

Categorical variables are measurements of individuals in terms of groups or categories where the gap between categories have no intrinsic meaning. A typical example of a categorical variable is race, where the gap between 'black' and 'white' has no proper interpretation, language, political affiliation, country of birth, and many other demographic variables.

It is vitally important to be able to distinguish between different data types because to a large extent these dictate what statistical techniques can be used. For example, it makes good sense to calculate the mean of a continuous variable but (as we have seen) no sense at all to calculate the mean of a categorical variable. The same idea extends to multivariate analysis. Some of the techniques we will look at work on correlation coefficients, which cannot be calculated for strictly categorical variables like race.

One further point on data types: some textbooks further divide numerical variables into *ratio-scaled* numerical variables and *interval-scaled* numerical variables; and divide categorical variables into *ordinal* categorical variables and *nominal* categorical variables. For the purposes of deciding which multivariate technique to use, this is an unnecessary detail and it is sufficient to know whether a variable is numerical or categorical. For the sake of completeness these additional terms are briefly described below. Ratio-scaled numerical variables are those that have a natural zero point (like age, height, and income). These are called "ratio-scaled" because they are not sensitive to units of measurement (if I am three times your height in meters I am also three times your height if it is measured in centimeters). This means that ratio-scaled variables have an arbitrary scale. Interval-scaled variables are still numeric but do not have a natural zero point (IQ, temperature in degrees Celcius, and most Likert-type rating scales are of this type). Interval-scaled variables therefore have an arbitrary zero point *and* an arbitrary scale. Ordinal categorical variables are those where the categories can be ordered even if the gaps between them cannot be interpreted (such as level of education, which can be ordered: none, primary-school, high-school, undergraduate degree, postgraduate degree). In contrast, the categories of a nominal categorical variable cannot be ordered in any meaningful way

(such as race or language group). It is also common to further classify numerical variables as *continuous* if they can take on any intermediate value on the scale (e.g. height) or *discrete* if the values a variable can take on are limited in some way (e.g. number of children).

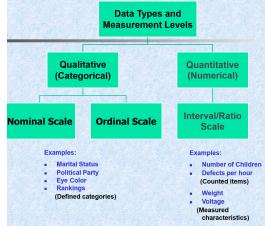


Figure 1: Data types

Dataset file formats

- Raw files, (.csv,.txt, .xlsx, .sav, etc.)
 - Databases (mySQL, MongoDB)
 - APIs (Twitter)
 - and others...
-

Types of EDA

There are broadly two types of EDA:

- Univariate
- Multivariate

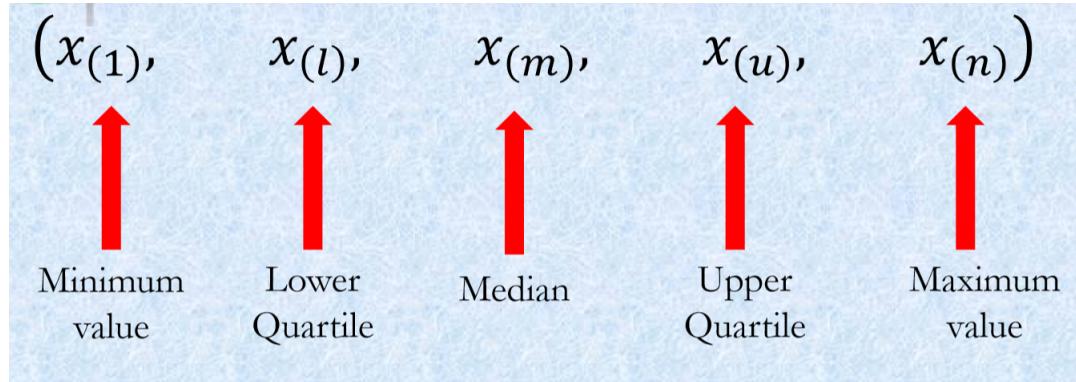
Most research in the business and social sciences makes use of some kind of multivariate analysis. Research that considers only one variable at a time (a *univariate* analysis) can provide useful information – for example, about the average rate of inflation over time, the variability of a particular share's return, or the relative proportion of the population that hold a certain opinion – but it is usually in the consideration of *relationships* between two or more variables that the most interesting and useful information is to be found. For example, what other variables are related to increases in the inflation rate or the rise in the price of

a particular share? Is it interest rates? Foreign exchange rates? And what causes people to prefer one opinion over another? Is it their education level? Income? The newspaper they read? Simply put, any analysis that considers the relationship between two or more variables is a *multivariate* analysis.

EDA methods - Summary Statistics

With summary statistics, our aim is to reduce a large data set to a few numbers which will help us understand the important features of the data.

- Compute a few “key” numbers: 5 number summaries.
- Let’s begin with the concept of ranked data.
- In a sample of size n , the smallest number has a rank of 1; the second smallest number has a rank of 2; ... ; the largest number has a rank of n .



$x_1, x_2, x_3, \dots, x_n$ and $x_{(r)}$ is the number with rank r .

- Frequency distributions
- Measures of central tendency and variability
- Smoothing techniques
- Analysis of tables
- Correlations

EDA methods - Plots

Tukey particularly emphasizes the value of graphs for discovery. Tukey’s approach to data analysis is highly visual and he has numerous suggestions for graphical displays. Tukey emphasizes the value of graphs for the following:

- Graphs can be used to store quantitative data,
- Graphs can be used to communicate conclusions,
- Graphs can be used to discover new information.

Some types of plots are better for one purpose, others are better for another.

Distribution of a Single Quantitative Variable

Tukey's novel distribution tools:

- Stem and Leaf: The measures Tukey proposes involve no arithmetic, only counting.
- Histogram: One should note,
 - * its height,
 - * where it is centered,
 - * how spread out it is,
 - * whether it is asymmetric,
 - * whether there are any discontinuities.

- Box-Whisker Plots: These plots show medians, quartiles, and two extreme values in a format that is easy to grasp quickly. Very powerful when comparing several frequency distributions.
 - Q-Q plots
-

Visual Display of a Single Qualitative Variable

- Frequency distribution
- Pie chart
- Bar chart

Visualising plays an important role in exploring your data, and you would know that Tukey favours analysis of data with four-color pen, graph paper, few tables etc.:




Figure 2: graph paper

Though we will use R and its functions for this purpose:

```
library(ggplot2)
library(tidyverse)
library(dplyr)
```

R Examples

The things that we are collecting data from (which could be people, shares, countries, animal species, songs ... anything you can collect data on) are called *cases* or *responses*. These appear as separate *rows* in the data matrix. The pieces of information that we use to describe each case are called *variables* or *attributes* and these appear in the *columns* of the data matrix. The *x*'s, remember, are simply placeholders for values to come. Specifically, the values to come may be numbers, or they may be words. It is perfectly allowable for the first column of *x*'s to be, for example, the first names of each person, e.g. $x_{11} = \text{Iris}$. Of course, this will affect the type of analysis we can do later on that variable (for example, it wouldn't make sense to calculate a mean' first name).

Example 1 - Female headed households in SA

[Womxn in Big Data South Africa: Female-Headed Households in South Africa competition](#)

The datasets are [provided](#) in a .csv file format, test.csv, train.csv, variable_descriptions.csv.

The *target* variable of interest is the percentage of households per ward that are both female-headed and earn an annual income that is below R19,600 (approximately \$2,300 USD in 2011).

```
Train <- read.csv("../Datasets/Woman/Train.csv", header = TRUE)
str(Train)
```

```
'data.frame': 2822 obs. of 63 variables:
 $ ward           : chr "41601001: Ward 1" "41601002: Ward 2" "41601003: Ward 3" "41601004: Ward 4" ...
 $ total_households: num 1674 1737 2404 1741 1731 ...
 $ total_individuals: num 5888 6735 7273 5734 6657 ...
 $ target          : num 16.8 21.5 10.9 23.1 13.7 ...
 $ dw_00           : num 0.934 0.697 0.811 0.66 0.951 ...
```

\$ dw_01	: num	0.000846 0.001253 0.004517 0 0.000655 ...
\$ dw_02	: num	0.00549 0.0044 0.00889 0.00613 0.00147 ...
\$ dw_03	: num	0.000676 0 0.003986 0 0.000598 ...
\$ dw_04	: num	0 0.002301 0.007735 0.000813 0.006999 ...
\$ dw_05	: num	0.001372 0.001323 0.000956 0.037245 0.000818 ...
\$ dw_06	: num	0.00575 0.00757 0.00669 0.00526 0.00498 ...
\$ dw_07	: num	0.03147 0.12355 0.02263 0.06891 0.00915 ...
\$ dw_08	: num	0.00808 0.15191 0.1299 0.21879 0.01538 ...
\$ dw_09	: num	0.00282 0.00149 0 0 0.00869 ...
\$ dw_10	: num	0.00143 0.00125 0 0 0 ...
\$ dw_11	: num	0.008224 0.00801 0.00415 0.002947 0.000673 ...
\$ dw_12	: int	0 0 0 0 0 0 0 0 0 0 ...
\$ dw_13	: int	0 0 0 0 0 0 0 0 0 0 ...
\$ psa_00	: num	0.26 0.29 0.186 0.281 0.197 ...
\$ psa_01	: num	0.608 0.55 0.677 0.593 0.518 ...
\$ psa_02	: num	0.000188 0 0.000489 0.000579 0.000989 ...
\$ psa_03	: num	0.01002 0.02134 0.02131 0.00725 0.00515 ...
\$ psa_04	: num	0.122 0.139 0.115 0.118 0.28 ...
\$ stv_00	: num	0.2835 0.1036 0.1658 0.0878 0.346 ...
\$ stv_01	: num	0.717 0.896 0.834 0.912 0.654 ...
\$ car_00	: num	0.274 0.145 0.272 0.128 0.405 ...
\$ car_01	: num	0.726 0.855 0.728 0.872 0.595 ...
\$ lln_00	: num	0.1188 0.0669 0.1 0.0292 0.1336 ...
\$ lln_01	: num	0.881 0.933 0.9 0.971 0.866 ...
\$ lan_00	: num	0.833 0.88 0.566 0.744 0.423 ...
\$ lan_01	: num	0.01234 0.00845 0.01599 0.00653 0.01435 ...
\$ lan_02	: num	0.001923 0.000328 0.001566 0.001188 0.000842 ...
\$ lan_03	: num	0.0509 0.0112 0.1113 0.0864 0.1219 ...
\$ lan_04	: num	0 0.000842 0.004795 0.006735 0.007027 ...
\$ lan_05	: num	0.000564 0.001759 0.002552 0.002308 0.002613 ...
\$ lan_06	: num	0.0761 0.0324 0.1481 0.1032 0.1474 ...
\$ lan_07	: num	0.00637 0.03084 0.13969 0.03828 0.08171 ...
\$ lan_08	: num	0.00366 0.00165 0.00317 0.00308 0.00304 ...
\$ lan_09	: num	0.000375 0.001308 0.000165 0.000582 0.000169 ...
\$ lan_10	: num	0.000372 0.000994 0.000779 0 0.000643 ...
\$ lan_11	: num	0.004943 0 0.001692 0.000197 0.001201 ...
\$ lan_12	: num	0.00272 0.00244 0.00251 0.00744 0.00428 ...
\$ lan_13	: int	0 0 0 0 0 0 0 0 0 0 ...
\$ lan_14	: num	0.006793 0.028061 0.0022 0.000174 0.192272 ...
\$ pg_00	: num	0.357 0.698 0.672 0.728 0.753 ...
\$ pg_01	: num	0.563 0.278 0.154 0.264 0.13 ...
\$ pg_02	: num	0.00426 0.0037 0.00218 0.00181 0.00452 ...
\$ pg_03	: num	0.072996 0.015835 0.167494 0.000956 0.106953 ...

```

$ pg_04          : num  0.00212 0.00404 0.00365 0.00539 0.00538 ...
$ lgt_00         : num  0.919 0.959 0.826 0.986 0.957 ...
$ pw_00          : num  0.743 0.309 0.323 0.677 0.771 ...
$ pw_01          : num  0.214 0.577 0.483 0.314 0.195 ...
$ pw_02          : num  0.01997 0.01895 0.08301 0.00269 0.0097 ...
$ pw_03          : num  0.00285 0.01457 0.05756 0 0.00486 ...
$ pw_04          : num  0.007537 0.057127 0.010358 0.000669 0.00129 ...
$ pw_05          : num  0 0.019092 0.001421 0 0.000673 ...
$ pw_06          : num  0.01293 0.00413 0.04088 0.00501 0.01763 ...
$ pw_07          : int   0 0 0 0 0 0 0 0 0 ...
$ pw_08          : int   0 0 0 0 0 0 0 0 0 ...
$ ADM4_PCODE     : chr  "ZA4161001" "ZA4161002" "ZA4161003" "ZA4161004" ...
$ lat            : num  -29.7 -29.1 -29.1 -29.4 -29.4 ...
$ lon            : num  24.7 24.8 25.1 24.9 25.3 ...
$ NL             : num  0.292 3.208 0 2.039 0 ...

```

```
summary(Train)
```

	ward	total_households	total_individuals	target
Length:	2822	Min. : 1	Min. : 402	Min. : 0.00
Class :	character	1st Qu.: 1779	1st Qu.: 7071	1st Qu.: 16.75
Mode :	character	Median : 2398	Median : 9367	Median : 24.16
		Mean : 3665	Mean : 12869	Mean : 24.51
		3rd Qu.: 3987	3rd Qu.: 14241	3rd Qu.: 32.23
		Max. :39685	Max. :91717	Max. :55.53
	dw_00	dw_01	dw_02	dw_03
	Min. :0.0000	Min. :0.000000	Min. :0.000000	Min. :0.0000000
	1st Qu.:0.5942	1st Qu.:0.002895	1st Qu.:0.002407	1st Qu.:0.0000000
	Median :0.7668	Median :0.010425	Median :0.005762	Median :0.0008066
	Mean :0.7122	Mean :0.092616	Mean :0.032043	Mean :0.0060567
	3rd Qu.:0.8817	3rd Qu.:0.068209	3rd Qu.:0.027913	3rd Qu.:0.0025383
	Max. :0.9950	Max. :0.931489	Max. :0.951806	Max. :0.2642393
	dw_04	dw_05	dw_06	dw_07
	Min. :0.0000000	Min. :0.0000000	Min. :0.000000	Min. :0.000000
	1st Qu.:0.0000000	1st Qu.:0.0000000	1st Qu.:0.002716	1st Qu.:0.004716
	Median :0.0006069	Median :0.0008654	Median :0.008639	Median :0.016295
	Mean :0.0086655	Mean :0.0062888	Mean :0.022374	Mean :0.039296
	3rd Qu.:0.0022246	3rd Qu.:0.0030272	3rd Qu.:0.025218	3rd Qu.:0.048731
	Max. :0.3920853	Max. :0.4359115	Max. :0.412936	Max. :0.455815
	dw_08	dw_09	dw_10	dw_11
	Min. :0.000000	Min. :0.0000000	Min. :0.0000000	Min. :0.000000
	1st Qu.:0.002888	1st Qu.:0.0002329	1st Qu.:0.0000000	1st Qu.:0.001991

Median :0.014991	Median :0.0017552	Median :0.0003909	Median :0.004092
Mean :0.064586	Mean :0.0068641	Mean :0.0011121	Mean :0.007902
3rd Qu.:0.074748	3rd Qu.:0.0065068	3rd Qu.:0.0010425	3rd Qu.:0.007803
Max. :0.798479	Max. :0.2828433	Max. :0.0687517	Max. :1.000000
dw_12	dw_13	psa_00	psa_01
Min. :0	Min. :0	Min. :0.0000	Min. :0.001293
1st Qu.:0	1st Qu.:0	1st Qu.:0.2556	1st Qu.:0.467217
Median :0	Median :0	Median :0.3017	Median :0.540874
Mean :0	Mean :0	Mean :0.3113	Mean :0.526568
3rd Qu.:0	3rd Qu.:0	3rd Qu.:0.3712	3rd Qu.:0.586087
Max. :0	Max. :0	Max. :0.5616	Max. :0.852493
psa_02	psa_03	psa_04	stv_00
Min. :0.0000000	Min. :0.00000	Min. :0.04279	Min. :0.0000
1st Qu.:0.0001326	1st Qu.:0.01698	1st Qu.:0.11014	1st Qu.:0.0982
Median :0.0003381	Median :0.02705	Median :0.12576	Median :0.1728
Mean :0.0005410	Mean :0.03369	Mean :0.12793	Mean :0.2259
3rd Qu.:0.0006835	3rd Qu.:0.04350	3rd Qu.:0.13973	3rd Qu.:0.3034
Max. :0.0194420	Max. :0.26738	Max. :0.99871	Max. :0.8405
stv_01	car_00	car_01	lln_00
Min. :0.1595	Min. :0.0000	Min. :0.04133	Min. :0.00000
1st Qu.:0.6966	1st Qu.:0.1310	1st Qu.:0.71851	1st Qu.:0.01732
Median :0.8272	Median :0.1780	Median :0.82197	Median :0.04014
Mean :0.7741	Mean :0.2503	Mean :0.74969	Mean :0.09764
3rd Qu.:0.9018	3rd Qu.:0.2815	3rd Qu.:0.86902	3rd Qu.:0.12087
Max. :1.0000	Max. :0.9587	Max. :1.00000	Max. :0.76261
lln_01	lan_00	lan_01	lan_02
Min. :0.2374	Min. :0.000000	Min. :0.000000	Min. :0.000000
1st Qu.:0.8791	1st Qu.:0.002842	1st Qu.:0.009433	1st Qu.:0.004081
Median :0.9599	Median :0.007914	Median :0.017589	Median :0.008956
Mean :0.9024	Mean :0.097603	Mean :0.058684	Mean :0.029416
3rd Qu.:0.9827	3rd Qu.:0.059327	3rd Qu.:0.036612	3rd Qu.:0.015081
Max. :1.0000	Max. :0.979246	Max. :0.939549	Max. :0.895365
lan_03	lan_04	lan_05	lan_06
Min. :0.000000	Min. :0.00000	Min. :0.000000	Min. :0.000000
1st Qu.:0.001647	1st Qu.:0.01034	1st Qu.:0.001675	1st Qu.:0.002681
Median :0.008835	Median :0.05253	Median :0.003986	Median :0.017154
Mean :0.039983	Mean :0.28432	Mean :0.116773	Mean :0.108053
3rd Qu.:0.039564	3rd Qu.:0.56850	3rd Qu.:0.055631	3rd Qu.:0.066745
Max. :0.852927	Max. :0.98616	Max. :0.978779	Max. :0.981207
lan_07	lan_08	lan_09	lan_10
Min. :0.000000	Min. :0.000000	Min. :0.0000000	Min. :0.0000000
1st Qu.:0.003906	1st Qu.:0.001675	1st Qu.:0.0002974	1st Qu.:0.0002999
Median :0.008403	Median :0.003045	Median :0.0012675	Median :0.0012002

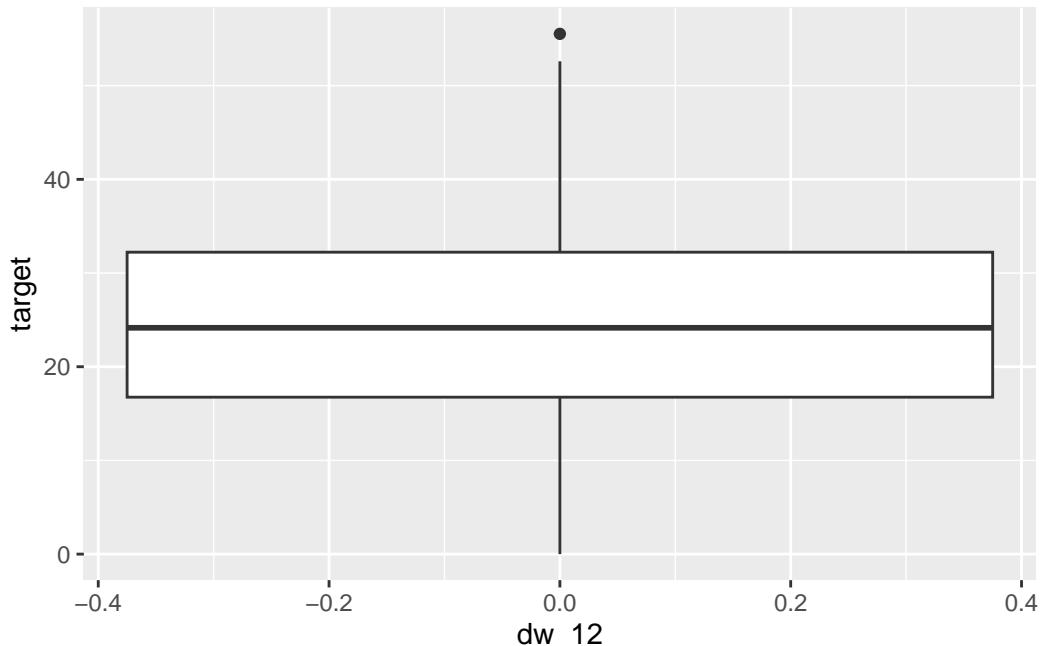
Mean : 0.130673	Mean : 0.004621	Mean : 0.0243186	Mean : 0.0242625
3rd Qu.: 0.065157	3rd Qu.: 0.005782	3rd Qu.: 0.0065378	3rd Qu.: 0.0052470
Max. : 0.963219	Max. : 0.034234	Max. : 0.9812332	Max. : 0.9828445
lan_11	lan_12	lan_13	lan_14
Min. : 0.000000	Min. : 0.000000	Min. : 0	Min. : 0.0000000
1st Qu.: 0.000495	1st Qu.: 0.002589	1st Qu.: 0	1st Qu.: 0.0000000
Median : 0.003261	Median : 0.006394	Median : 0	Median : 0.0001459
Mean : 0.053985	Mean : 0.012809	Mean : 0	Mean : 0.0145029
3rd Qu.: 0.029783	3rd Qu.: 0.013722	3rd Qu.: 0	3rd Qu.: 0.0121078
Max. : 0.991674	Max. : 0.367785	Max. : 0	Max. : 0.9984484
pg_00	pg_01	pg_02	pg_03
Min. : 0.01105	Min. : 0.000000	Min. : 0.0000000	Min. : 0.0000000
1st Qu.: 0.87528	1st Qu.: 0.001015	1st Qu.: 0.0008769	1st Qu.: 0.0004514
Median : 0.98975	Median : 0.003124	Median : 0.0017966	Median : 0.0012081
Mean : 0.86214	Mean : 0.040938	Mean : 0.0187979	Mean : 0.0744293
3rd Qu.: 0.99562	3rd Qu.: 0.012582	3rd Qu.: 0.0048827	3rd Qu.: 0.0418406
Max. : 1.00000	Max. : 0.969519	Max. : 0.9395640	Max. : 0.9405628
pg_04	lgt_00	pw_00	pw_01
Min. : 0.0000000	Min. : 0.001692	Min. : 0.000000	Min. : 0.000000
1st Qu.: 0.0006644	1st Qu.: 0.796471	1st Qu.: 0.08764	1st Qu.: 0.1113
Median : 0.0016958	Median : 0.914061	Median : 0.27800	Median : 0.3021
Mean : 0.0036926	Mean : 0.836432	Mean : 0.35969	Mean : 0.3297
3rd Qu.: 0.0041264	3rd Qu.: 0.964334	3rd Qu.: 0.58295	3rd Qu.: 0.5088
Max. : 0.3678423	Max. : 1.000000	Max. : 0.99591	Max. : 0.9376
pw_02	pw_03	pw_04	pw_05
Min. : 0.000000	Min. : 0.000000	Min. : 0.0000000	Min. : 0.0000000
1st Qu.: 0.008673	1st Qu.: 0.002099	1st Qu.: 0.0007147	1st Qu.: 0.0001595
Median : 0.069065	Median : 0.016496	Median : 0.0051637	Median : 0.0014590
Mean : 0.127555	Mean : 0.041589	Mean : 0.0196551	Mean : 0.0110081
3rd Qu.: 0.183384	3rd Qu.: 0.058626	3rd Qu.: 0.0250545	3rd Qu.: 0.0094322
Max. : 1.000000	Max. : 0.327393	Max. : 0.3067867	Max. : 0.2282606
pw_06	pw_07	pw_08	ADM4_PCODE
Min. : 0.000000	Min. : 0	Min. : 0	Length: 2822
1st Qu.: 0.005217	1st Qu.: 0	1st Qu.: 0	Class : character
Median : 0.025165	Median : 0	Median : 0	Mode : character
Mean : 0.110818	Mean : 0	Mean : 0	Median : -26.55
3rd Qu.: 0.116638	3rd Qu.: 0	3rd Qu.: 0	Mean : -26.88
Max. : 0.961522	Max. : 0	Max. : 0	3rd Qu.: -25.57
lon	NL		lat
Min. : 16.76	Min. : 0.000		Min. : -32.49
1st Qu.: 27.71	1st Qu.: 3.033		1st Qu.: -28.57
Median : 28.96	Median : 9.206		
Mean : 28.67	Mean : 17.438		Max. : -22.33

```
3rd Qu.:30.44    3rd Qu.:26.891  
Max.     :32.86    Max.     :63.000
```

Now we will explore this dataset.

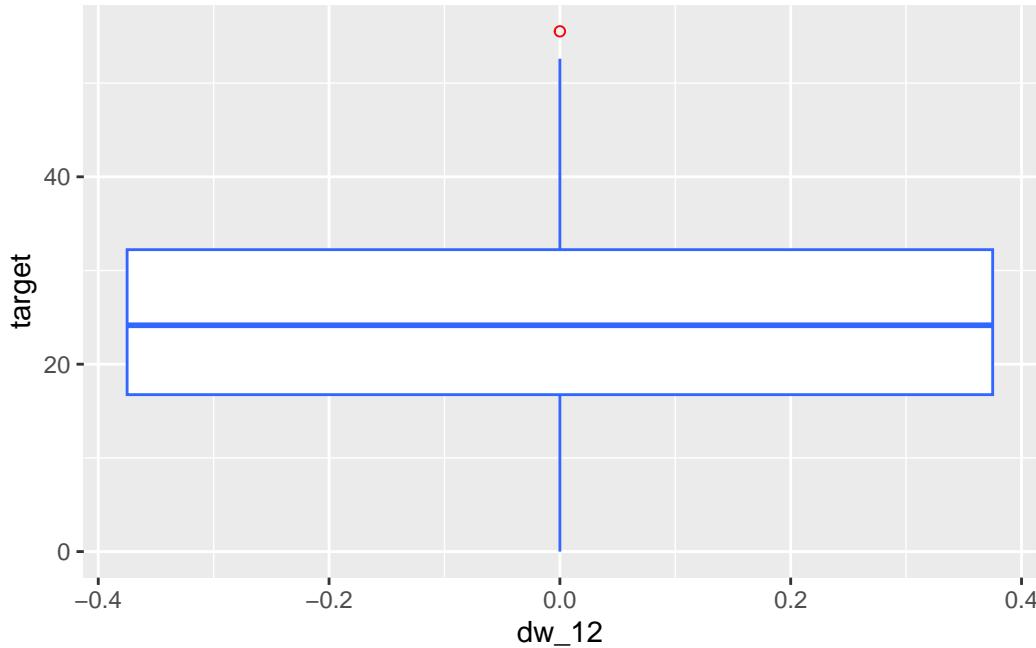
The majority of the packages that you will use are part of the so-called tidyverse package:

```
#install.packages("tidyverse")  
library(tidyverse)  
  
plot1 = ggplot(data = Train, aes(dw_12,target))  
plot1 + geom_boxplot()
```



How to interpret this: IntroStat p:21-26

```
plot1 + geom_boxplot(fill = "white", colour = "#3366FF", outlier.colour = "red", outlier.shape = 19)
```



💡 Tip

[ggplot cheatsheet](#)

[ggplot2: Elegant Graphics for Data Analysis \(Use R!\) by Hadley Wickham](#)

Example 2 - Gapminder

```
#install.packages("gapminder")
# install.packages("pacman")
# p_load(ggplot2, gapminder)
library(pacman)
p_load(ggplot2, gapminder, tidyverse) # This will first check if the package is installed,
gapminder::gapminder # this is to use the function without loading it.
```

```
# A tibble: 1,704 x 6
```

```

country continent year lifeExp      pop gdpPercap
<fct>     <fct>   <int>    <dbl>    <int>    <dbl>
1 Afghanistan Asia     1952     28.8  8425333     779.
2 Afghanistan Asia     1957     30.3  9240934     821.
3 Afghanistan Asia     1962     32.0 10267083     853.
4 Afghanistan Asia     1967     34.0 11537966     836.
5 Afghanistan Asia     1972     36.1 13079460     740.
6 Afghanistan Asia     1977     38.4 14880372     786.
7 Afghanistan Asia     1982     39.9 12881816     978.
8 Afghanistan Asia     1987     40.8 13867957     852.
9 Afghanistan Asia     1992     41.7 16317921     649.
10 Afghanistan Asia    1997     41.8 22227415     635.
# i 1,694 more rows

```

```
gapminder
```

```

# A tibble: 1,704 x 6
  country continent year lifeExp      pop gdpPercap
  <fct>     <fct>   <int>    <dbl>    <int>    <dbl>
  1 Afghanistan Asia     1952     28.8  8425333     779.
  2 Afghanistan Asia     1957     30.3  9240934     821.
  3 Afghanistan Asia     1962     32.0 10267083     853.
  4 Afghanistan Asia     1967     34.0 11537966     836.
  5 Afghanistan Asia     1972     36.1 13079460     740.
  6 Afghanistan Asia     1977     38.4 14880372     786.
  7 Afghanistan Asia     1982     39.9 12881816     978.
  8 Afghanistan Asia     1987     40.8 13867957     852.
  9 Afghanistan Asia     1992     41.7 16317921     649.
  10 Afghanistan Asia    1997     41.8 22227415     635.
# i 1,694 more rows

```

Frequency Distribution

```

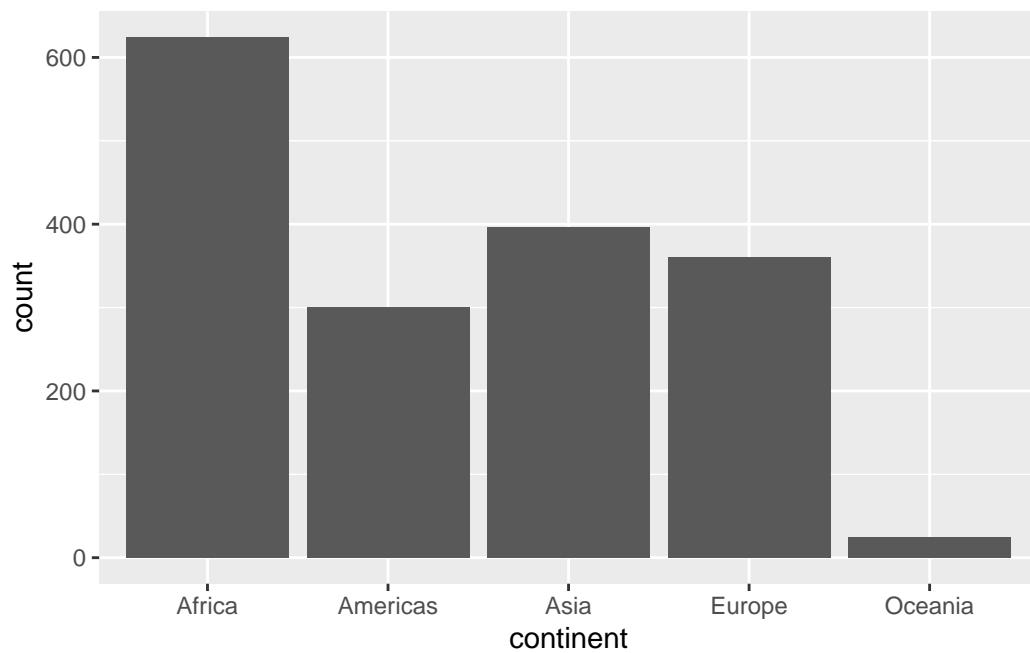
library(ggplot2)
table(gapminder$continent)

```

Africa	Americas	Asia	Europe	Oceania
624	300	396	360	24

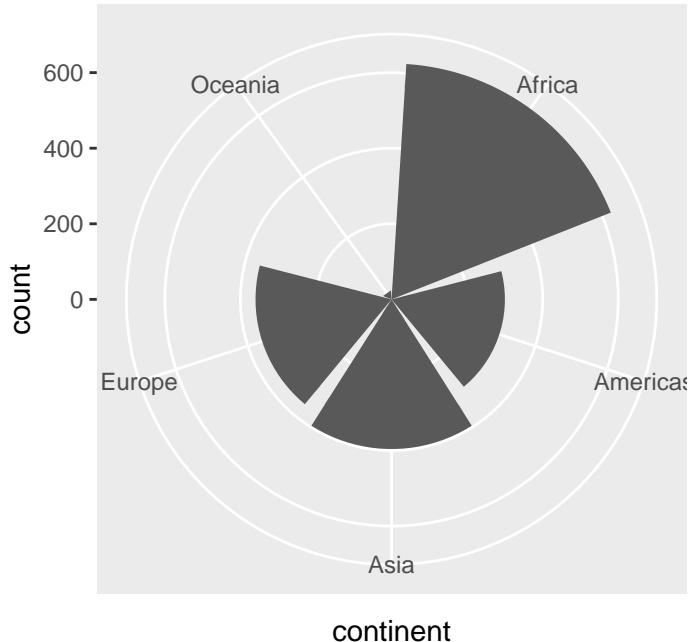
Bar Plot

```
library(ggplot2)
plot1 <- ggplot(gapminder, aes(x=continent)) + geom_bar()
plot1
```



Pie Chart

```
plot1 + coord_polar()
```

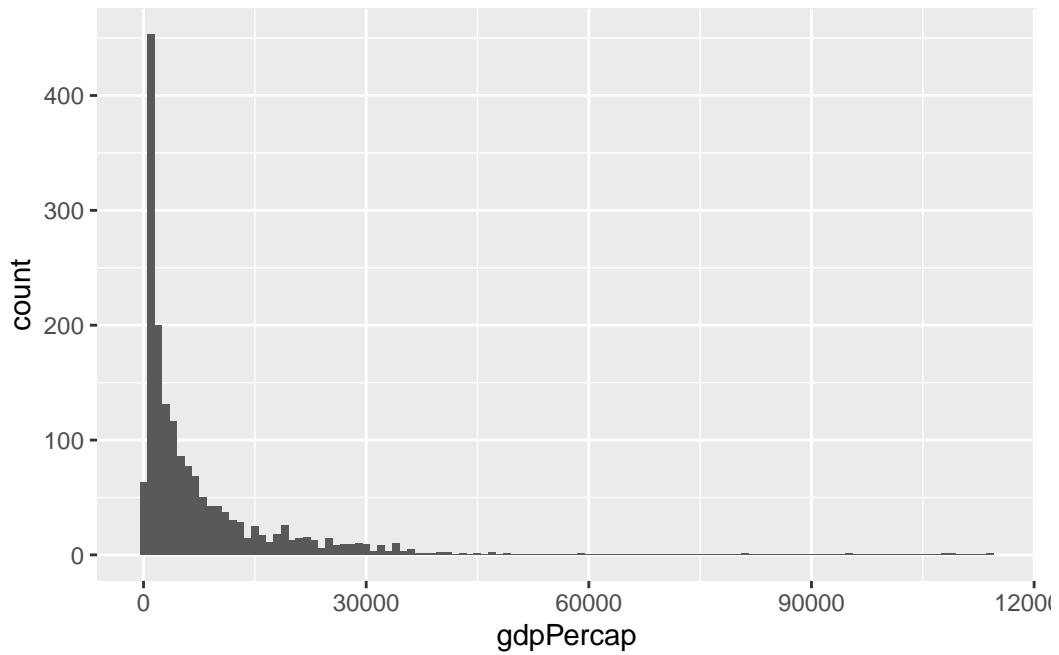


If you would like to have a regular pie chart, then you need to provide the frequency distribution.

Histogram

A Simple Histogram

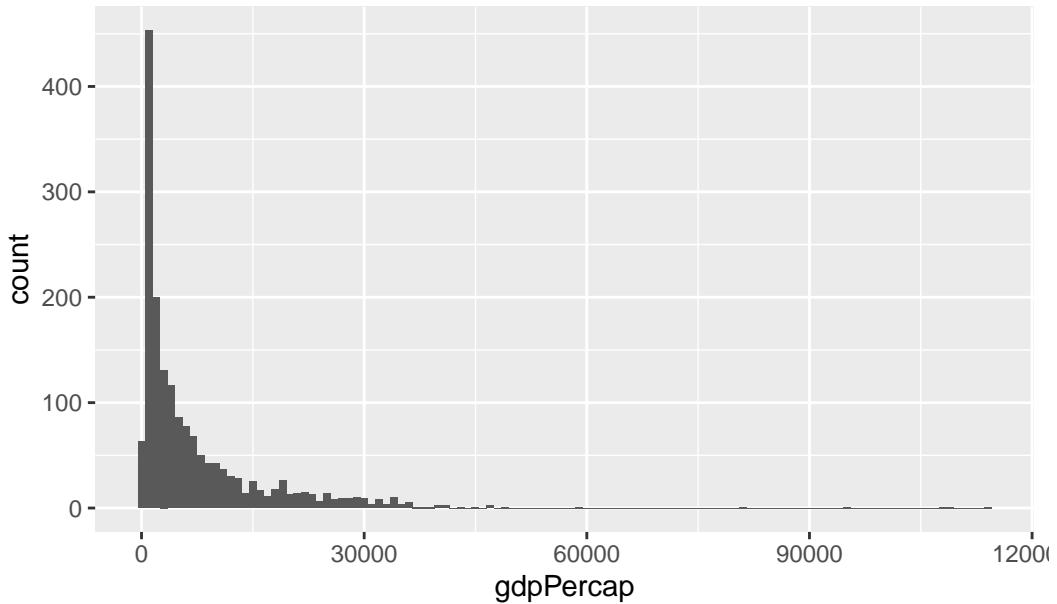
```
library(ggplot2)
plot2 <- ggplot(gapminder,
                 aes(x = gdpPercap))
plot2 + geom_histogram(binwidth = 1000)
```



Histogram With a Title

```
plot2 +  
  geom_histogram(binwidth = 1000) +  
  labs(title = "Histogram of GDP per Capita for All Years")
```

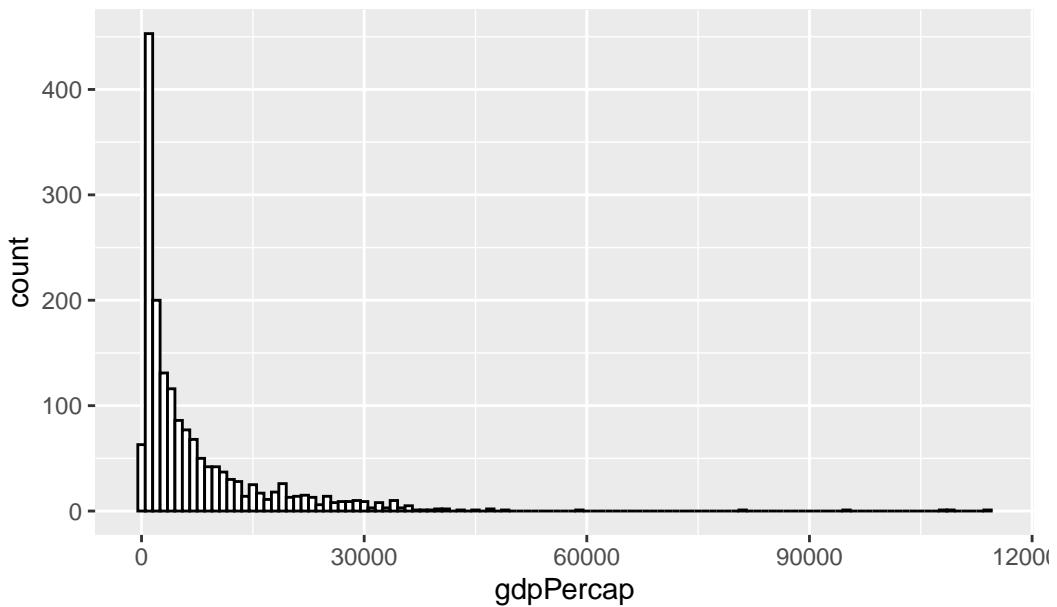
Histogram of GDP per Capita for All Years



Histogram with Different Color Schemes:

```
plot2 +  
  geom_histogram(binwidth = 1000, color="black", fill="white") +  
  labs(title = "Histogram of GDP per Capita for All Years")
```

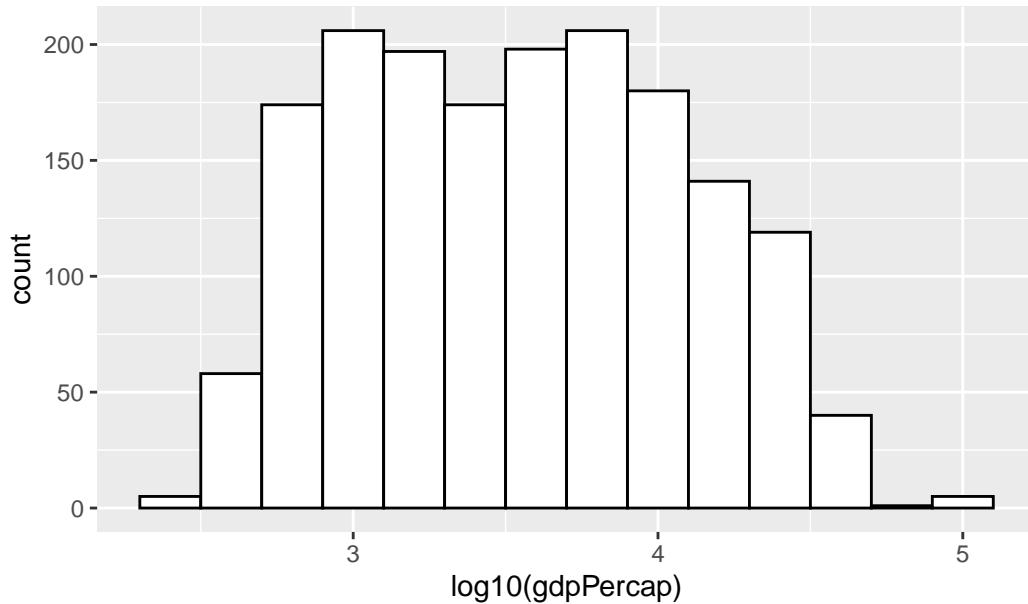
Histogram of GDP per Capita for All Years



Histogram of Log Transformed Variable:

```
plot3 <- ggplot(gapminder,
                 aes(x = log10(gdpPercap)))
plot3 +
  geom_histogram(binwidth = .2, color="black", fill="white") +
  labs(title = "Histogram of Log Transformed GDP per Capita for All Years")
```

Histogram of Log Transformed GDP per Capita for All Years



Determine the Binwidth

How do we determine the binwidth?

- Sturges' rule uses class intervals of length

$$L = \frac{x_{\max} - x_{\min}}{1 + 1.44 \ln(n)}$$

- Genstat rule uses class intervals of length:


$$L = \frac{x_{\max} - x_{\min}}{\sqrt{n}}$$

- or a general rule

So we can create our own function for the binwidth:

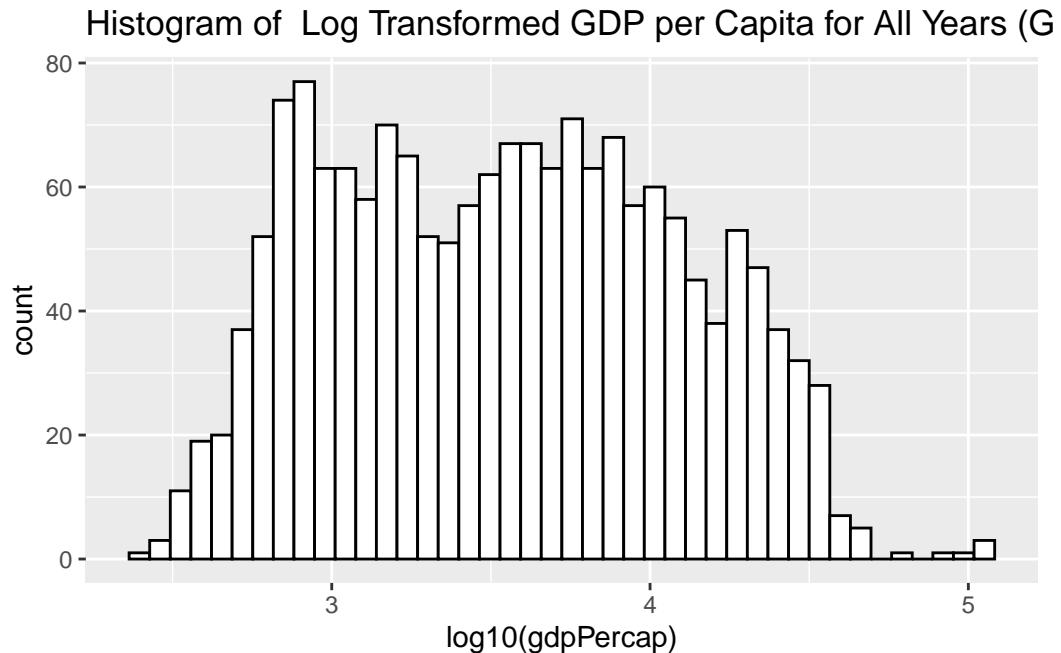
```
width_bin = function(x) (max(x)-min(x)) / sqrt(length(x))
manualbin = width_bin(log10(gapminder$gdpPerCap))
```

```
plot3 +
  geom_histogram(binwidth = manualbin, color="black", fill="white")
```



or simply

```
plot3 +
  geom_histogram(binwidth = function(x) (max(x)-min(x)) / sqrt(length(x)), color="black", fill="white",
  labs(title = "Histogram of Log Transformed GDP per Capita for All Years (Genstat Binwid")
```




But you will notice that Gdp per capita variable includes all years, all continents, all countries!!!

Histogram for a Subset of Data

Log Transformed GDP per Capita for 2007:

```
plot4 <- ggplot(subset(gapminder, year == 2007),
                 aes(x = log10(gdpPerCap)))
plot4 +
  geom_histogram(binwidth = function(x) (max(x)-min(x)) / sqrt(length(x)), color="black",
                 fill="white")
```

Histogram of Log Transformed GDP per Capita for 2007 (Gens)

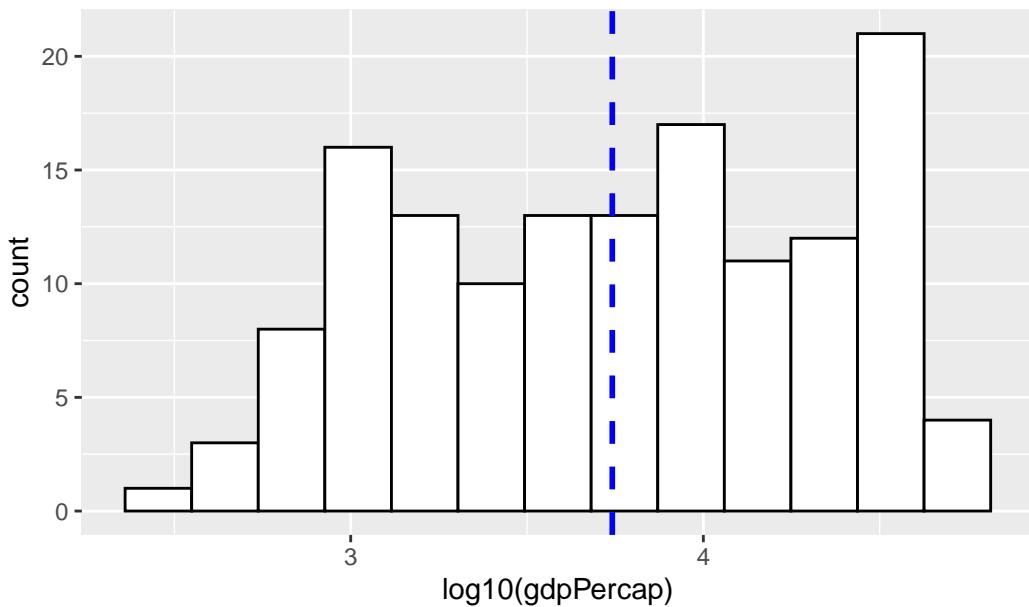


Histogram with Overall Mean Line

Log Transformed GDP per Capita for 2007 with the Overall Mean Line


```
# Histogram with mean of log10(gdpPerCap) on the plot
plot4 +
  geom_histogram(binwidth = function(x) (max(x)-min(x)) / sqrt(length(x)), color="black", fill="white") +
  geom_vline(aes(xintercept=mean(log10(gdpPerCap))), color="blue", linetype="dashed", linewidth=1) +
  labs(title = "Histogram of Log Transformed GDP per Capita for 2007 (Genstat Binwidth)")
```

Histogram of Log Transformed GDP per Capita for 2007 (Gens)



Histogram with Density plot

```
# Histogram with density plot
ggplot(subset(gapminder, year == 2007),
       aes(x = log10(gdpPerCap))) +
  geom_histogram(aes(y=after_stat(density)), binwidth = function(x) (max(x)-min(x)) / sqrt(12))
  geom_density(alpha=0, fill="#FF6666") #alpha for transparency, if alpha = 0, no fill
```




Histogram with Facets

How about looking at the differences among different continents?

```
# Histogram with mean of log10(gdpPerCap) on the plot
plot4 +
  geom_histogram(binwidth = function(x) (max(x)-min(x)) / sqrt(length(x)), color="black", fill="white") +
  geom_vline(aes(xintercept=mean(log10(gdpPerCap))),
             color="blue", linetype="dashed", size=1) +
  facet_grid(continent ~ .)
```

Warning: Using `size` aesthetic for lines was deprecated in ggplot2 3.4.0.
i Please use `linewidth` instead.



Boxplots

```
# Histogram with mean of log10(gdpPercap) on the plot
plot5 <- ggplot(subset(gapminder, year == 2007),
                 aes(x = year, y = log10(gdpPercap)))
# if x axis variable is numeric, then one single boxplot
# if x axis variable is categorical, then works like facets

plot5 +
  geom_boxplot() #+ coord_flip()
```




Try with “continent” variable.

Scatter Plots

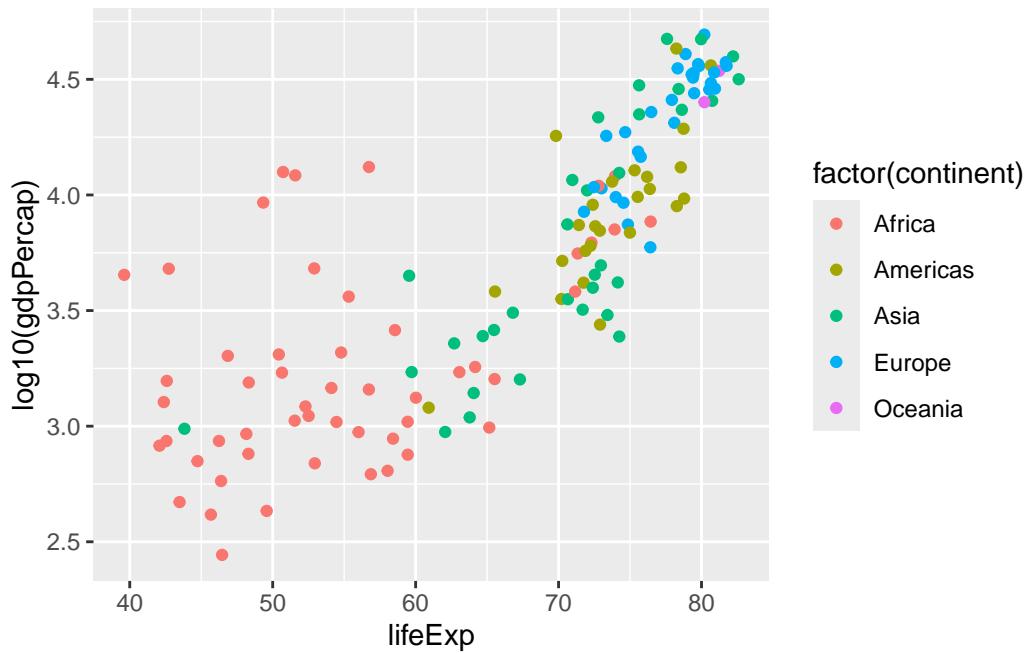
A Simple Scatter Plot

```
plot6 <- ggplot(subset(gapminder, year == 2007),  
                 aes(x = lifeExp, y = log10(gdpPercap)))  
plot6 +  
  geom_point()
```



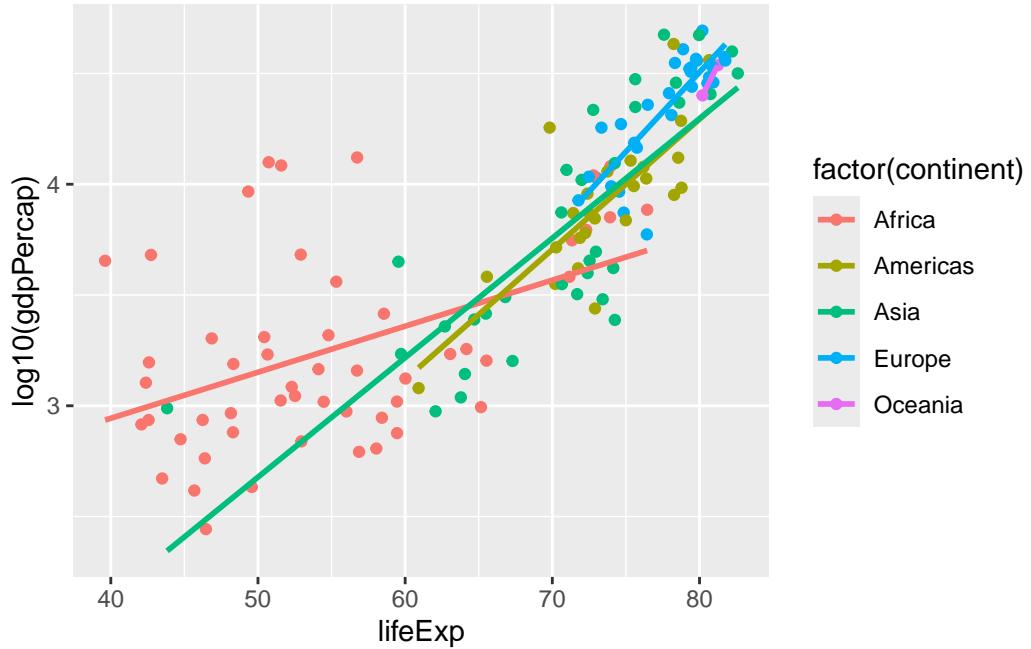
Scatter Plot with Labellings

```
plot6 +  
  geom_point(aes(colour = factor(continent)))
```




Scatter Plot with Linear Lines for Different Groups

```
plot6 +
  geom_point(aes(colour = factor(continent))) +
  geom_smooth(aes(group = continent, colour = factor(continent)), lwd = 1, se = FALSE, method =
`geom_smooth()` using formula = 'y ~ x'
```



```
plot7 <- ggplot(gapminder,
                  aes(x = lifeExp, y = log10(gdpPercap)))
plot7 +
  geom_point(aes(colour = factor(continent))) +
  facet_wrap(~ year) # scales = "free_x"
```



Class exercise

[Thesis 1](#) [Thesis 2](#) [Thesis 3](#), [Figure 7.1](#) [Thesis 4](#) [Thesis 5](#) [Thesis 6 - Figures 7-9](#) [Thesis 7](#) [Thesis 8](#) [Thesis 9](#) [Thesis 10](#) [Thesis 11](#)