

STA5092Z EDA Lecture 5-6

Table of contents

EDA Lecture 5-6 - Checking your data - Checklist	1
Workflow - scripts/projects	1
Libraries	2
Boston Dataset	2
Ames Dataset	10
Happiness Dataset	10
Class Exercises	10
Pokemon Dataset	10

EDA Lecture 5-6 - Checking your data - Checklist

- Formulate your question
- Check your data
- Automate your project workflow

Workflow - scripts/projects

- Do not let R to auto save your workspace, go to options in Rstudio, set “save workspace to .RData on exit” option to “never”. This will enforce you to record every step you take while dealing with data and analysis.
- Create an RStudio project for each data analysis project.
- Keep data files there.
- Write your R scripts in an R markdown file. Keep scripts there.
- Start your script with the packages that you need.
- Do not include install.packages code in a script you share.
- Save your outputs (plots and cleaned data) there.

- Only ever use relative paths, not absolute paths. Because they hinder sharing: no one else will have exactly the same directory configuration as you.
- DO NOT PRINT EVERYTHING!!!
- R Studio projects keep all your files in a specified folder, data, output, Rmd or R files, etc.
- When you shut down R, you can easily go to your specified folder and open .Rproj file and all your work will be together.
- For graphics, you can print your plots in separate pdf files with ggsave() function and all will be saved in your project folder.

We will go through these using the Boston, Ames, World Happiness Datasets. Before moving on, we will look at the code chunk options:

- **echo**: Show (TRUE) or hide (FALSE) the source code in output,
- **eval**: Execute (TRUE) or skip (FALSE) the code,
- **include**: Include both code and output in the rendered document (FALSE hides both),
- **warning**: Show (TRUE) or suppress (FALSE) warnings,
- **message**: Show (TRUE) or suppress (FALSE) messages,
- **results**: “markup” (default), “asis”, “hide”, “hold”. Controls how results are displayed,
- **cache**: Cache results for faster re-rendering,
- **error**: Show errors in output (TRUE) or stop rendering (FALSE).

Libraries

```
Loading required package: pacman
```

Boston Dataset

1) Read in your data:

2) Run str() or glimpse()

```
'data.frame': 506 obs. of 13 variables:
 $ crim    : num  0.00632 0.02731 0.02729 0.03237 0.06905 ...
 $ zn      : num  18 0 0 0 0 12.5 12.5 12.5 12.5 ...
 $ indus   : num  2.31 7.07 7.07 2.18 2.18 2.18 7.87 7.87 7.87 ...
 $ chas    : int  0 0 0 0 0 0 0 0 0 ...
 $ nox     : num  0.538 0.469 0.469 0.458 0.458 0.458 0.524 0.524 0.524 ...
 $ rm      : num  6.58 6.42 7.18 7 7.15 ...
 $ age     : num  65.2 78.9 61.1 45.8 54.2 58.7 66.6 96.1 100 85.9 ...
 $ dis     : num  4.09 4.97 4.97 6.06 6.06 ...
```

```

$ rad      : int  1 2 2 3 3 3 5 5 5 5 ...
$ tax      : int 296 242 242 222 222 222 311 311 311 311 ...
$ ptratio: num 15.3 17.8 17.8 18.7 18.7 18.7 15.2 15.2 15.2 15.2 ...
$ lstat   : num 4.98 9.14 4.03 2.94 5.33 ...
$ medv    : num 24 21.6 34.7 33.4 36.2 28.7 22.9 27.1 16.5 18.9 ...

```

Rows: 506

Columns: 13

```

$ crim    <dbl> 0.00632, 0.02731, 0.02729, 0.03237, 0.06905, 0.02985, 0.08829, ~
$ zn       <dbl> 18.0, 0.0, 0.0, 0.0, 0.0, 0.0, 12.5, 12.5, 12.5, 12.5, 1~
$ indus   <dbl> 2.31, 7.07, 7.07, 2.18, 2.18, 2.18, 7.87, 7.87, 7.87, 7.~
$ chas    <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ~
$ nox     <dbl> 0.538, 0.469, 0.469, 0.458, 0.458, 0.458, 0.524, 0.524, 0.524, ~
$ rm      <dbl> 6.575, 6.421, 7.185, 6.998, 7.147, 6.430, 6.012, 6.172, 5.631, ~
$ age     <dbl> 65.2, 78.9, 61.1, 45.8, 54.2, 58.7, 66.6, 96.1, 100.0, 85.9, 9~
$ dis     <dbl> 4.0900, 4.9671, 4.9671, 6.0622, 6.0622, 6.0622, 5.5605, 5.9505~
$ rad     <int> 1, 2, 2, 3, 3, 3, 5, 5, 5, 5, 5, 5, 4, 4, 4, 4, 4, 4, 4, 4, ~
$ tax     <int> 296, 242, 242, 222, 222, 311, 311, 311, 311, 311, 311, 31~
$ ptratio <dbl> 15.3, 17.8, 17.8, 18.7, 18.7, 18.7, 15.2, 15.2, 15.2, 15.2, 15~
$ lstat   <dbl> 4.98, 9.14, 4.03, 2.94, 5.33, 5.21, 12.43, 19.15, 29.93, 17.10~
$ medv    <dbl> 24.0, 21.6, 34.7, 33.4, 36.2, 28.7, 22.9, 27.1, 16.5, 18.9, 15~
```

3) Look at the top and the bottom of your data, does that make sense? Any rows appearing as NAs by mistake?

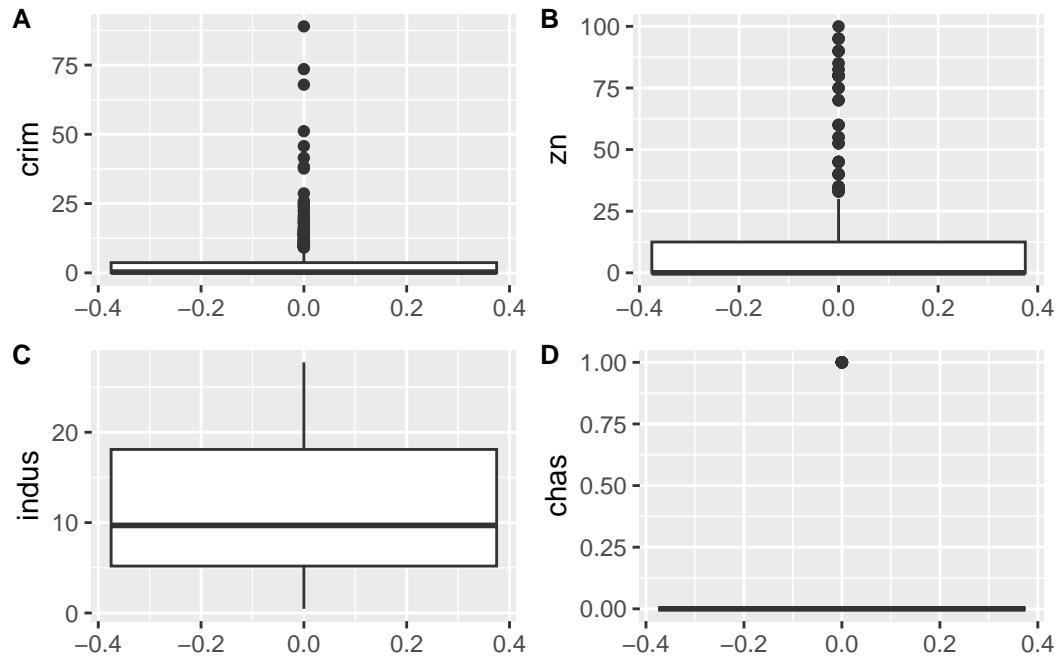
	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	lstat	medv
1	0.00632	18	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	4.98	24.0
2	0.02731	0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	9.14	21.6
3	0.02729	0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	4.03	34.7

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	lstat	medv
504	0.06076	0	11.93	0	0.573	6.976	91.0	2.1675	1	273	21	5.64	23.9
505	0.10959	0	11.93	0	0.573	6.794	89.3	2.3889	1	273	21	6.48	22.0
506	0.04741	0	11.93	0	0.573	6.030	80.8	2.5050	1	273	21	7.88	11.9

4) Did you check the total number of observations? n

[1] 506 13

5) Did you plot univariate summaries of the data (histograms, density plots, boxplots)?



```
Saving 5.5 x 3.5 in image
```

6) Did you consider correlations between variables (scatterplots)?

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax
crim	1.000	-0.200	0.407	-0.056	0.421	-0.219	0.353	-0.380	0.626	0.583
zn	-0.200	1.000	-0.534	-0.043	-0.517	0.312	-0.570	0.664	-0.312	-0.315
indus	0.407	-0.534	1.000	0.063	0.764	-0.392	0.645	-0.708	0.595	0.721
chas	-0.056	-0.043	0.063	1.000	0.091	0.091	0.087	-0.099	-0.007	-0.036
nox	0.421	-0.517	0.764	0.091	1.000	-0.302	0.731	-0.769	0.611	0.668
rm	-0.219	0.312	-0.392	0.091	-0.302	1.000	-0.240	0.205	-0.210	-0.292
age	0.353	-0.570	0.645	0.087	0.731	-0.240	1.000	-0.748	0.456	0.506
dis	-0.380	0.664	-0.708	-0.099	-0.769	0.205	-0.748	1.000	-0.495	-0.534

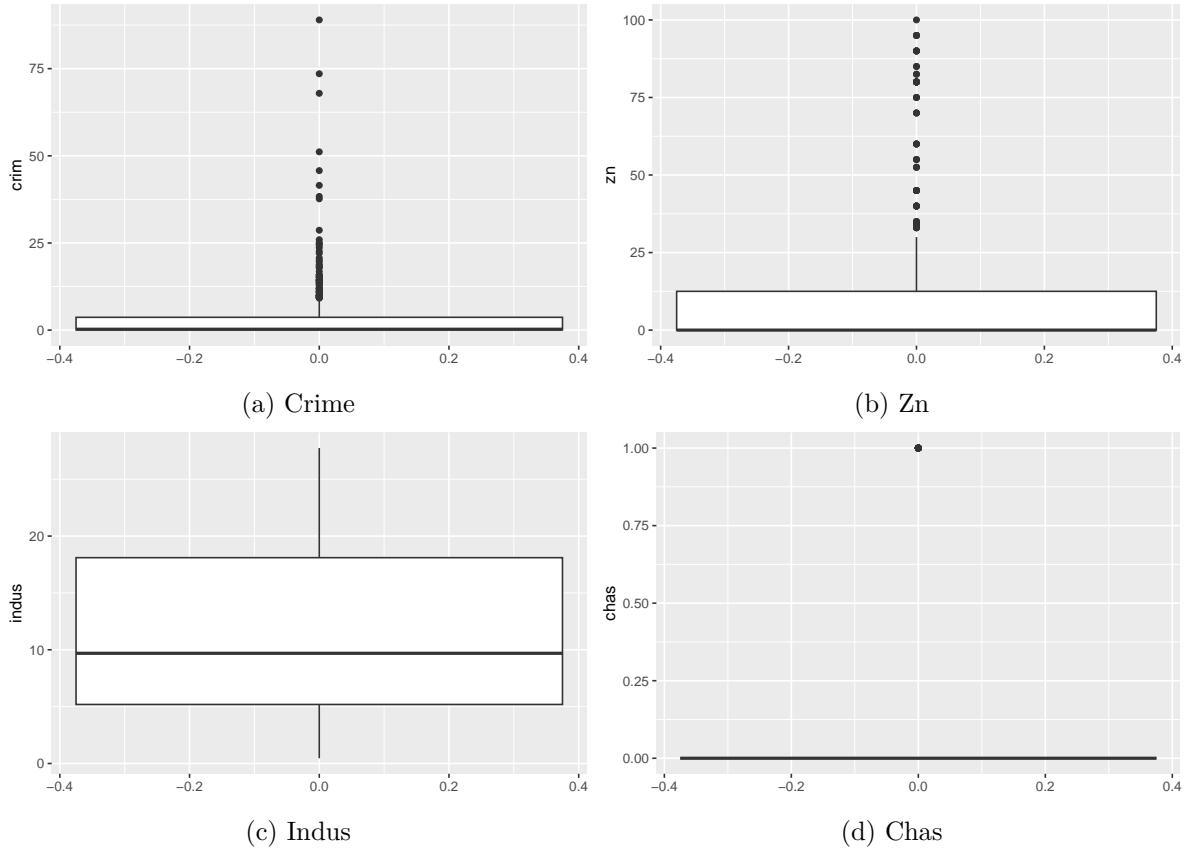
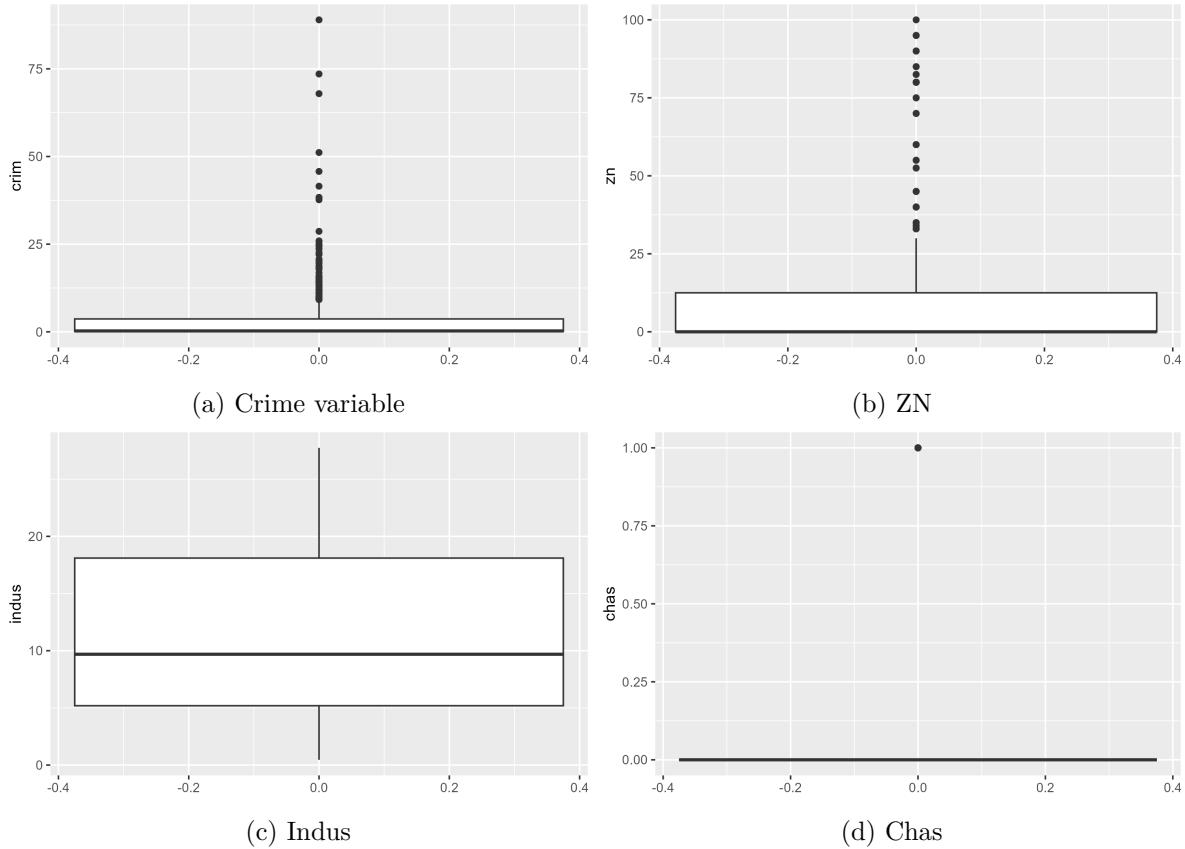
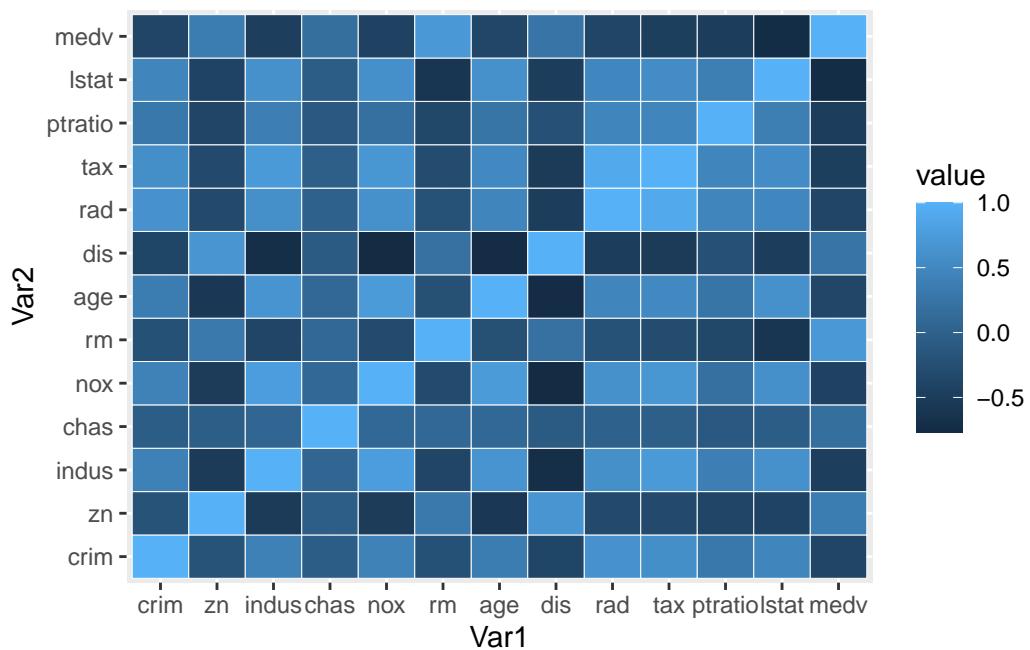
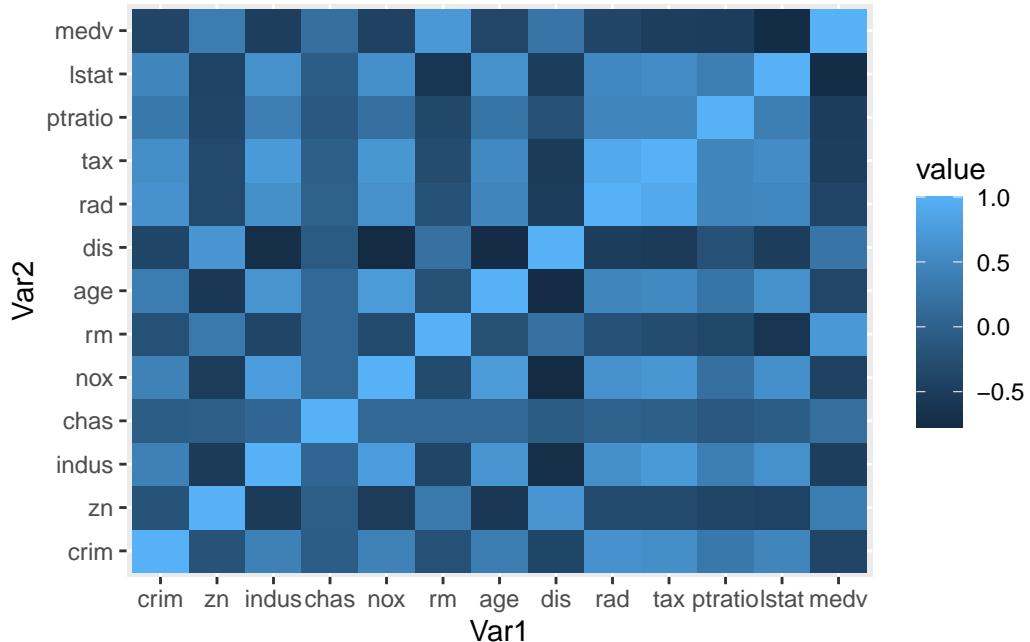


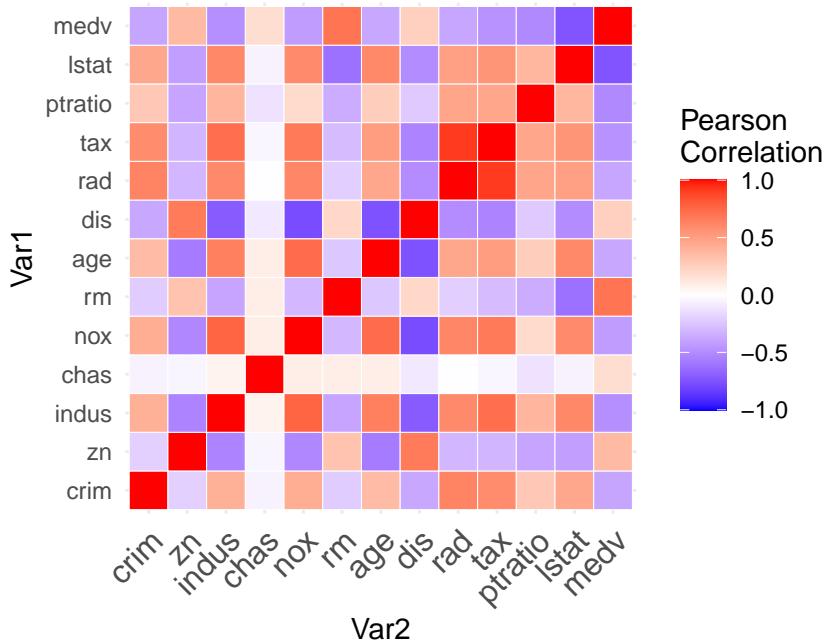
Figure 1: Model diagnostics



rad	0.626	-0.312	0.595	-0.007	0.611	-0.210	0.456	-0.495	1.000	0.910
tax	0.583	-0.315	0.721	-0.036	0.668	-0.292	0.506	-0.534	0.910	1.000
ptratio	0.290	-0.392	0.383	-0.122	0.189	-0.356	0.262	-0.232	0.465	0.461
lstat	0.456	-0.413	0.604	-0.054	0.591	-0.614	0.602	-0.497	0.489	0.544
medv	-0.388	0.360	-0.484	0.175	-0.427	0.695	-0.377	0.250	-0.382	-0.469
	ptratio	lstat	medv							
crim	0.290	0.456	-0.388							
zn	-0.392	-0.413	0.360							
indus	0.383	0.604	-0.484							
chas	-0.122	-0.054	0.175							
nox	0.189	0.591	-0.427							
rm	-0.356	-0.614	0.695							
age	0.262	0.602	-0.377							
dis	-0.232	-0.497	0.250							
rad	0.465	0.489	-0.382							
tax	0.461	0.544	-0.469							
ptratio	1.000	0.374	-0.508							
lstat	0.374	1.000	-0.738							
medv	-0.508	-0.738	1.000							

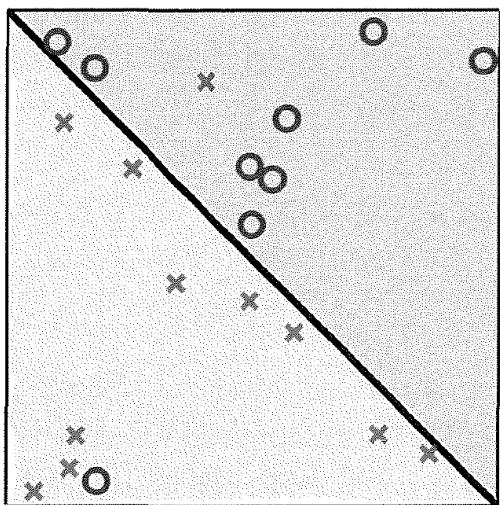
	Var1	Var2	value
1	crim	crim	1.000
2	zn	crim	-0.200
3	indus	crim	0.407
4	chas	crim	-0.056
5	nox	crim	0.421
6	rm	crim	-0.219



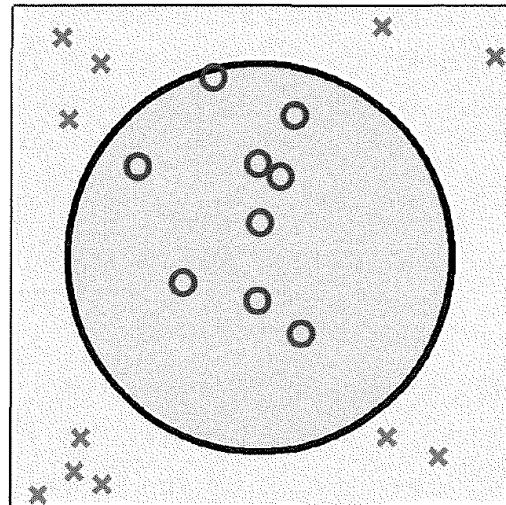


7) Did you check for outliers? What is an outlier? Single

variable outlier or multivariate outlier? Why so important? What do we do? Throw away the samples? Man made sampling bias! ##### 8) Did you identify missing values? ##### 9) Did you check the units of all data points to make sure they are in the right range? ##### 10) Did you try to identify any errors or miscoding of variables? ##### 11) Did you consider plotting on a log scale or any other transformation? Use log transforms for ratio measurements. ##### 12) For large data sets, subsample before plotting ##### 13) Use color and size to check for confounding



(a) Few noisy data.



(b) Nonlinearly separable.

Figure 3.1: Data sets that are not linearly separable but are (a) linearly separable after discarding a few examples, or (b) separable by a more sophisticated curve.

Figure 3: Noisy data plots (Ref: Learning from data, A short course by Yaser Abu-Mostafa et al, p.79)

[**Ames Dataset**](#)

[**Happiness Dataset**](#)

Class Exercises

Pokemon Dataset

Exercise 1

Select the first three columns of the pokemon dataset using their column names.

Exercise 2

Select all the columns of the pokemon dataset except “abilities”.

Exercise 3

Select all columns of the pokemon dataset that start with the character string “against”.

Exercise 4

Filter the rows of the pokemon dataset for $\text{against_bug} \geq 1$ and $\text{against_dark} \geq 1$.

Exercise 5

Arrange rows by a particular column, such as the against_bug .

Exercise 6

Select three columns from pokemon, arrange the rows by against_bug , then arrange the rows by against_dark .

Exercise 7

Create a new column called proportion, which is the ratio of height to weight of the pokemon.

Exercise 8

Compute the average damage that a pokemon can do against bug, apply the `mean()` function to the column against_bug , and call the summary value “ave_against_bug”. HINT: Use `summarize()`.

Exercise 10

Obtain the same summary statistics for different types (Type1) of pokemon. HINT: Use `group_by()`.