

Unsupervised Learning Methods

Dr Sebnem Er

2020-10-12

Contents

1	Introduction	5
2	Association Rules	7
2.1	Prerequisites	7
2.2	The Groceries Dataset	7
2.3	Support Count (Item Frequencies) and Item Frequency Plot . . .	10
2.4	Support	11
2.5	Rule Generation with Apriori Algorithm	13
2.6	Using your own dataset stored as a csv file	19
2.7	References:	20
3	Cluster Analysis	21
3.1	Prerequisites	21
3.2	K-means clustering	22
3.3	K-medoids clustering	23
3.4	Hierarchical Clustering	24
3.5	Methods for determining number of clusters	28
3.6	Clustering with CLARA	36
3.7	Clustering with DBSCAN	36
3.8	Clustering using mixture models	38
3.9	Cluster Profiling	40
4	Self Organising Maps	43

Chapter 1

Introduction

This book will guide you through the R codes for the following Unsupervised Learning methods:

- Association Rules
- Cluster Analysis
- Self Organising Maps

The chapters will be made available on Tuesdays when we start a new topic. So please update your browser to access the codes for the relevant chapter.

Chapter 2

Association Rules

2.1 Prerequisites

You need to have the following R packages installed and recalled into your library:

```
library(datasets)
library(grid)
library(tidyverse)
library(readxl)
library(knitr)
library(ggplot2)
library(lubridate)
library(arules)
library(arulesViz)
library(plyr)
```

2.2 The Groceries Dataset

We shall mine Groceries dataset for association rules using the Apriori Algorithm. The Groceries dataset can be loaded from R. The steps for doing so are shown below. Note that you will only be able to load the data set once the package `arules` has been loaded into R. The Groceries dataset contains a collection of receipts with each line representing 1 receipt and the items purchased. Each line is called a transaction and each column in a row represents an item.

```

data(Groceries)
summary(Groceries)

## transactions as itemMatrix in sparse format with
## 9835 rows (elements/itemsets/transactions) and
## 169 columns (items) and a density of 0.02609146
##
## most frequent items:
##      whole milk other vegetables      rolls/buns      soda
##           2513           1903           1809           1715
##           yogurt      (Other)
##           1372           34055
##
## element (itemset/transaction) length distribution:
## sizes
##      1      2      3      4      5      6      7      8      9     10     11     12     13     14     15     16
## 2159 1643 1299 1005  855  645  545  438  350  246  182  117  78   77   55   46
##      17     18     19     20     21     22     23     24     26     27     28     29     32
##      29     14     14      9     11      4      6      1      1      1      1      3      1
##
##      Min. 1st Qu.  Median      Mean 3rd Qu.      Max.
##      1.000  2.000   3.000   4.409   6.000  32.000
##
## includes extended item information - examples:
##      labels level2      level1
## 1 frankfurter sausage meat and sausage
## 2      sausage sausage meat and sausage
## 3 liver loaf sausage meat and sausage

```

As you can see, the data is in “transactions” format with a density of 0.0261 (check slides to remember what this value means). There are 9835 transactions with 169 distinct items that can be bought in this database (D).

The summary function also provides the distribution of number items per transaction and the most popular items.

Now let us examine the first 3 transactions in D .

```

inspect(head(Groceries, 3))

##      items
## [1] {citrus fruit,
##      semi-finished bread,
##      margarine,
##      ready soups}

```



```
## [2] {tropical fruit,
##      yogurt,
##      coffee}
## [3] {whole milk}
```

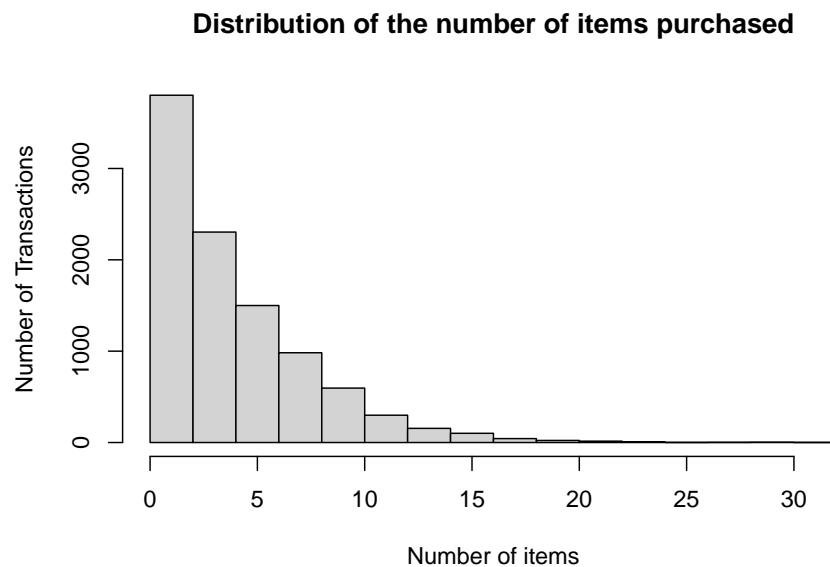
The first customer bought {citrus fruit,semi-finished bread,margarine,ready soups}, whereas the third customer bought only {whole milk}.

We can also find how many items each transaction contains, for the first 10 transactions:

```
head(size(Groceries), 10)
```

```
## [1] 4 3 1 4 4 5 1 5 1 2
```

```
hist(size(Groceries), main = "Distribution of the number of items purchased", xlab = "Number of i
```



As it is clear, the distribution of the number of items is skewed to right, clearly most transactions include fewer number of items, only very few have more than 10 items purchased together.

2.3 Support Count (Item Frequencies) and Item Frequency Plot

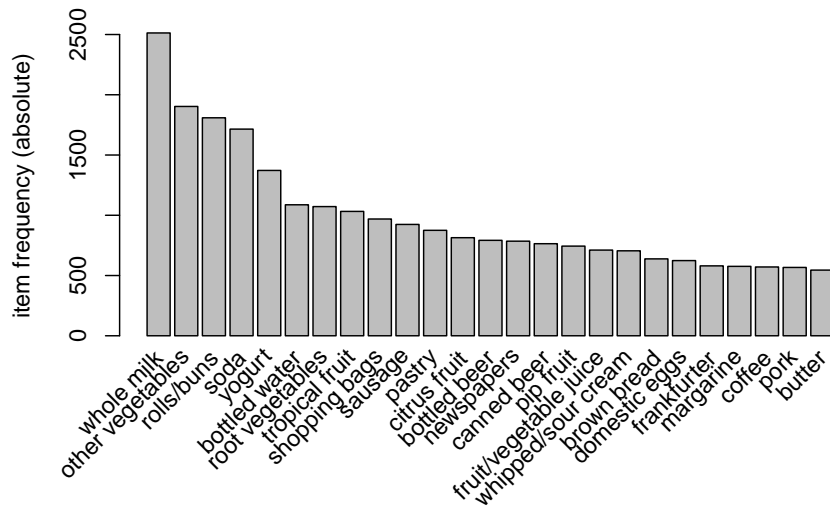
We can check the support count ($freq(A)$) for the top 25 products with the following R code:

```
itemSupportCount = itemFrequency(Groceries, type = "absolute") # obtain the counts for
itemSupportCount = sort(itemSupportCount, decreasing = TRUE) # sort the counts in desc
head(itemSupportCount, 25) # check the support count for the top 25 items
```

##	whole milk	other vegetables	rolls/buns
##	2513	1903	1809
##	soda	yogurt	bottled water
##	1715	1372	1087
##	root vegetables	tropical fruit	shopping bags
##	1072	1032	969
##	sausage	pastry	citrus fruit
##	924	875	814
##	bottled beer	newspapers	canned beer
##	792	785	764
##	pip fruit	fruit/vegetable juice	whipped/sour cream
##	744	711	705
##	brown bread	domestic eggs	frankfurter
##	638	624	580
##	margarine	coffee	pork
##	576	571	567
##	butter		
##	545		

We can also plot the support count, it is possible to change the colours of the bars as well.

```
itemFrequencyPlot(Groceries, topN = 25, type="absolute")
```



We can see that top purchased product is {whole milk} and it appears in 2513 transactions out of 9835. Therefore the support count for {whole milk} is 2513.

2.4 Support

Remember the support ($S(A)$) is calculated as follows:

$$S(A) = \frac{\text{freq}(A)}{n}$$

The support for {whole milk} would be

$$S(\text{whole milk}) = \frac{\text{freq}(\text{whole milk})}{n} = \frac{2513}{9835} = 25.55\%$$

It is possible to obtain this information with the same code as shown previously by simply replacing `type="absolute"` with the `type="relative"` option:


```
itemSupport = itemFrequency(Groceries, type = "relative") # obtain the counts for individual items
itemSupport = sort(itemSupport, decreasing = TRUE) # sort the counts in descending order
head(itemSupport, 25) # check the support count for the top 25 items
```

##	whole milk	other vegetables	rolls/buns
##	0.25551601	0.19349263	0.18393493

##	soda	yogurt	bottled water
##	0.17437722	0.13950178	0.11052364
##	root vegetables	tropical fruit	shopping bags
##	0.10899847	0.10493137	0.09852567
##	sausage	pastry	citrus fruit
##	0.09395018	0.08896797	0.08276563
##	bottled beer	newspapers	canned beer
##	0.08052872	0.07981698	0.07768175
##	pip fruit	fruit/vegetable juice	whipped/sour cream
##	0.07564820	0.07229283	0.07168277
##	brown bread	domestic eggs	frankfurter
##	0.06487036	0.06344687	0.05897306
##	margarine	coffee	pork
##	0.05856634	0.05805796	0.05765125
##	butter		
##	0.05541434		

We can also plot the support.

```
itemFrequencyPlot(Groceries, topN = 25, type="relative")
```



Note that the maximum support is low. To ensure that the top 25 frequent items are included in the analysis the minimum support would have to be less than 0.10! (10%) Suppose we set the minimum support to 0.001 and minimum confidence to 0.8. We can mine some rules by executing the following R code:

2.5 Rule Generation with Apriori Algorithm

- We are going to use the Apriori algorithm within the `arules` library to mine frequent itemsets and association rules..
- Assume that we want to generate all the rules that satisfy the support threshold of 0.1% and confidence threshold of 80%, then we need to enter `supp=0.001` and `conf=0.8` values in the `apriori()` function. If you want stronger rules, you can increase the value of `conf` and for more extended rules give higher value to `maxlen`.
- It might be desirable to sort the rules according either confidence or support, here we chose sorting according to confidence in a descending manner.
- Finally we can examine the rules using `summary()` function.

```
rules <- apriori(Groceries, parameter = list(supp=0.001, conf=0.8))

## Apriori
##
## Parameter specification:
## confidence minval  smax  arem  aval originalSupport  maxtime support minlen
##           0.8     0.1    1 none   FALSE             TRUE       5   0.001    1
## maxlen target  ext
##          10   rules TRUE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose
##       0.1 TRUE TRUE  FALSE TRUE    2    TRUE
##
## Absolute minimum support count: 9
##
## set item appearances ...[0 item(s)] done [0.00s].
## set transactions ...[169 item(s), 9835 transaction(s)] done [0.01s].
## sorting and recoding items ... [157 item(s)] done [0.00s].
## creating transaction tree ... done [0.00s].
## checking subsets of size 1 2 3 4 5 6 done [0.04s].
## writing ... [410 rule(s)] done [0.06s].
## creating S4 object ... done [0.00s].

rules <- sort(rules, by='confidence', decreasing = TRUE)
summary(rules)

## set of 410 rules
```

```
##
## rule length distribution (lhs + rhs):sizes
##   3   4   5   6
## 29 229 140  12
##
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##    3.000   4.000   4.000   4.329   5.000   6.000
##
## summary of quality measures:
##      support      confidence      coverage      lift
## Min.      :0.001017   Min.      :0.8000   Min.      :0.001017   Min.      : 3.131
## 1st Qu.:0.001017   1st Qu.:0.8333   1st Qu.:0.001220   1st Qu.: 3.312
## Median :0.001220   Median :0.8462   Median :0.001322   Median : 3.588
## Mean    :0.001247   Mean    :0.8663   Mean    :0.001449   Mean    : 3.951
## 3rd Qu.:0.001322   3rd Qu.:0.9091   3rd Qu.:0.001627   3rd Qu.: 4.341
## Max.    :0.003152   Max.    :1.0000   Max.    :0.003559   Max.    :11.235
##      count
## Min.      :10.00
## 1st Qu.:10.00
## Median :12.00
## Mean    :12.27
## 3rd Qu.:13.00
## Max.    :31.00
##
## mining info:
##      data ntransactions support confidence
## Groceries          9835   0.001         0.8
```

In this output we are provided with the following information:

- There are 410 rules based on 0.001 support and 0.8 confidence thresholds.
- The distribution of the number of items in each rule (rule length distribution): Most rules are 4 items long.

We need use the `inspect()` function to see the actual rules.

```
inspect(rules[1:5])
```

```
##      lhs                                rhs      support confidence      coverage      lift
## [1] {rice,                                => {whole milk} 0.001220132      1 0.001220132 3.91364
##      sugar}
## [2] {canned fish,                        => {whole milk} 0.001118454      1 0.001118454 3.91364
##      hygiene articles}
## [3] {root vegetables,
```

```
##      butter,
##      rice}          => {whole milk} 0.001016777      1 0.001016777 3.913649      10
## [4] {root vegetables,
##      whipped/sour cream,
##      flour}          => {whole milk} 0.001728521      1 0.001728521 3.913649      17
## [5] {butter,
##      soft cheese,
##      domestic eggs}  => {whole milk} 0.001016777      1 0.001016777 3.913649      10
```

If we look at the confidence we see that for the top 5 rules it is 1, this indicates 100% confidence:


- 100% customers who bought “{rice, sugar}” end up buying “{whole milk}” as well.
- 100% customers who bought “{canned fish, hygiene articles}” end up buying “{whole milk}” as well.

In the following section we will look at visualizing the rules.

2.5.1 Visualisation of the Rules

```
topRules <- rules[1:10]
plot(topRules)
```


```
## To reduce overplotting, jitter is added! Use jitter = 0 to prevent jitter.
```



The scatter plot of support and confidence of the top ten rules shows us that high confidence rules have low support values.

```
plot(rules)
```

```
## To reduce overplotting, jitter is added! Use jitter = 0 to prevent jitter.
```

In the following section we will look at removing redundant rules.

2.5.2 Removing redundant rules

You may want to remove rules that are subsets of larger rules. Use the code below to remove such rules:

```
subset.rules <- which(colSums(is.subset(rules, rules)) > 1) # get subset rules in vector
# is.subset() determines if elements of one vector contain all the elements of other
length(subset.rules)
```

```
## [1] 91
```

```
subset.rules <- rules[-subset.rules] # remove subset rules.
```

2.5.3 Finding rules related to given items

In the case of specific product in interest, either as a precedent (LHS) or as a consequent (RHS) in the rule, you need to set the “**appearance=**” parameter in the apriori rule generating function:

Let us say we are interested in those transactions that end up in buying “root vegetables”:

```
rveg.rules <- apriori(Groceries, parameter = list(supp=0.001, conf=0.8), appearance = 1)
```

```
## Apriori
##
## Parameter specification:
## confidence minval smax arem aval originalSupport maxtime support minlen
##           0.8    0.1    1 none FALSE                TRUE      5    0.001      1
## maxlen target  ext
##      10  rules TRUE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose
##    0.1 TRUE TRUE  FALSE TRUE    2    TRUE
##
## Absolute minimum support count: 9
##
## set item appearances ...[1 item(s)] done [0.00s].
## set transactions ...[169 item(s), 9835 transaction(s)] done [0.00s].
## sorting and recoding items ... [157 item(s)] done [0.00s].
## creating transaction tree ... done [0.00s].
## checking subsets of size 1 2 3 4 5 6 done [0.02s].
## writing ... [5 rule(s)] done [0.00s].
## creating S4 object ... done [0.00s].
```

```
inspect(head(rveg.rules))
```

	lhs	rhs	support	confidence	coverage
## [1]	{other vegetables, whole milk, yogurt, rice}	=> {root vegetables}	0.001321810	0.8666667	0.001525165
## [2]	{tropical fruit, other vegetables, whole milk, oil}	=> {root vegetables}	0.001321810	0.8666667	0.001525165
## [3]	{beef, citrus fruit, tropical fruit, other vegetables}	=> {root vegetables}	0.001016777	0.8333333	0.001220132
## [4]	{citrus fruit, other vegetables, soda, fruit/vegetable juice}	=> {root vegetables}	0.001016777	0.9090909	0.001118454
## [5]	{tropical fruit,				

```
##      other vegetables,
##      whole milk,
##      yogurt,
##      oil}                      => {root vegetables} 0.001016777 0.9090909 0.001118454 8.340400
```

2.6 Using your own dataset stored as a csv file

You might want to use a dataset from a csv file. The format of this file should be as follows:

- Transactions in the rows (remember in our small example, we had 5 transactions.)
- Items per transaction should be entered separately in different columns (items were A, B, C, D, E, and F)

	A	B	C	D	E
1	A	D			
2	A	B	E	C	
3	D	B	F	C	
4	A	D	B	C	
5	A	D	B	F	
6					
7					
8					
9					
10					

Figure 2.1: How the data looks like in csv format:

- The data should be extracted using the `read.transactions()` function.

```
slideExample <- read.transactions('C:/Users/01438475/Google Drive/UCTcourses/Analytics/Unsupervised Learning/transactions.csv')
slideExample
```

```
## transactions in sparse format with
```

```
## 5 transactions (rows) and  
## 6 items (columns)
```

```
inspect(head(slideExample, 6))
```

```
##      items  
## [1] {A,D}  
## [2] {A,B,C,E}  
## [3] {B,C,D,F}  
## [4] {A,B,C,D}  
## [5] {A,B,D,F}
```

```
size(head(slideExample))
```

```
## [1] 2 4 4 4 4
```

I will leave all the rest for you to obtain.

2.7 References:

- R and Data Mining
- Susan Li - MBA
- Datacamp
- Dr Juwa Nyirenda's lecture notes

Chapter 3

Cluster Analysis

We will use the built-in R dataset `USArrests` which contains statistics, in arrests per 100,000 residents for assault, murder, and rape in each of the 50 US states in 1973. It includes also the percent of the population living in urban areas. It contains 50 observations on 4 variables:

3.1 Prerequisites

We will need the following packages:

```
library(cluster)
library(NbClust)
library(fpc)
```

Load the data set

```
data("USArrests")
# Remove any missing value (i.e, NA values for not available)
# That might be present in the data
df <- na.omit(USArrests)
# View the first 6 rows of the data
head(df, n = 6)
```

##	Murder	Assault	UrbanPop	Rape
## Alabama	13.2	236	58	21.2
## Alaska	10.0	263	48	44.5
## Arizona	8.1	294	80	31.0
## Arkansas	8.8	190	50	19.5

```
## California    9.0    276    91 40.6
## Colorado     7.9    204    78 38.7
```

Before clustering is done, we can compute some descriptive statistics for the data

```
desc_stats <- data.frame(
  Min = apply(df, 2, min), # minimum
  Med = apply(df, 2, median), # median
  Mean = apply(df, 2, mean), # mean
  SD = apply(df, 2, sd), # Standard deviation
  Max = apply(df, 2, max) # Maximum
)
desc_stats <- round(desc_stats, 1)
head(desc_stats)
```

```
##           Min   Med  Mean   SD   Max
## Murder    0.8   7.2   7.8  4.4  17.4
## Assault  45.0 159.0 170.8 83.3 337.0
## UrbanPop 32.0  66.0  65.5 14.5  91.0
## Rape      7.3  20.1  21.2  9.4  46.0
```

Note that the variables have large different means and variances. Therefore we need to standardise them.

```
df <- scale(df)
head(df)
```

```
##           Murder  Assault  UrbanPop  Rape
## Alabama  1.24256408 0.7828393 -0.5209066 -0.003416473
## Alaska   0.50786248 1.1068225 -1.2117642  2.484202941
## Arizona   0.07163341 1.4788032  0.9989801  1.042878388
## Arkansas  0.23234938 0.2308680 -1.0735927 -0.184916602
## California 0.27826823 1.2628144  1.7589234  2.067820292
## Colorado  0.02571456 0.3988593  0.8608085  1.864967207
```

For partition clustering methods we will assume that $K=2$ clusters

3.2 K-means clustering

We will use the `kmeans()` function in the `stats` package.

```
set.seed(123)
km.out <- kmeans(df, 2, nstart = 25)
# k-means group number of each observation
km.out$cluster
```

```
##      Alabama      Alaska      Arizona      Arkansas      California
##      1          1          1          2          1
##      Colorado  Connecticut  Delaware      Florida      Georgia
##      1          2          2          1          1
##      Hawaii      Idaho      Illinois      Indiana      Iowa
##      2          2          1          2          2
##      Kansas      Kentucky  Louisiana      Maine      Maryland
##      2          2          1          2          1
##      Massachusetts  Michigan  Minnesota  Mississippi  Missouri
##      2          1          2          1          1
##      Montana      Nebraska      Nevada  New Hampshire  New Jersey
##      2          2          1          2          2
##      New Mexico      New York  North Carolina  North Dakota      Ohio
##      1          1          1          2          2
##      Oklahoma      Oregon      Pennsylvania  Rhode Island  South Carolina
##      2          2          2          2          1
##      South Dakota      Tennessee      Texas          Utah      Vermont
##      2          1          1          2          2
##      Virginia      Washington  West Virginia  Wisconsin      Wyoming
##      2          2          2          2          2
```

3.3 K-medoids clustering

We will use the `pam()` in the `cluster` package.

```
pam.out <- pam(df, 2)
pam.out$cluster
```

```
##      Alabama      Alaska      Arizona      Arkansas      California
##      1          1          1          2          1
##      Colorado  Connecticut  Delaware      Florida      Georgia
##      1          2          2          1          1
##      Hawaii      Idaho      Illinois      Indiana      Iowa
##      2          2          1          2          2
##      Kansas      Kentucky  Louisiana      Maine      Maryland
##      2          2          1          2          1
##      Massachusetts  Michigan  Minnesota  Mississippi  Missouri
##      2          1          2          1          1
```

##	Montana	Nebraska	Nevada	New Hampshire	New Jersey
##	2	2	1	2	2
##	New Mexico	New York	North Carolina	North Dakota	Ohio
##	1	1	1	2	2
##	Oklahoma	Oregon	Pennsylvania	Rhode Island	South Carolina
##	2	2	2	2	1
##	South Dakota	Tennessee	Texas	Utah	Vermont
##	2	1	1	2	2
##	Virginia	Washington	West Virginia	Wisconsin	Wyoming
##	2	2	2	2	2

3.4 Hierarchical Clustering

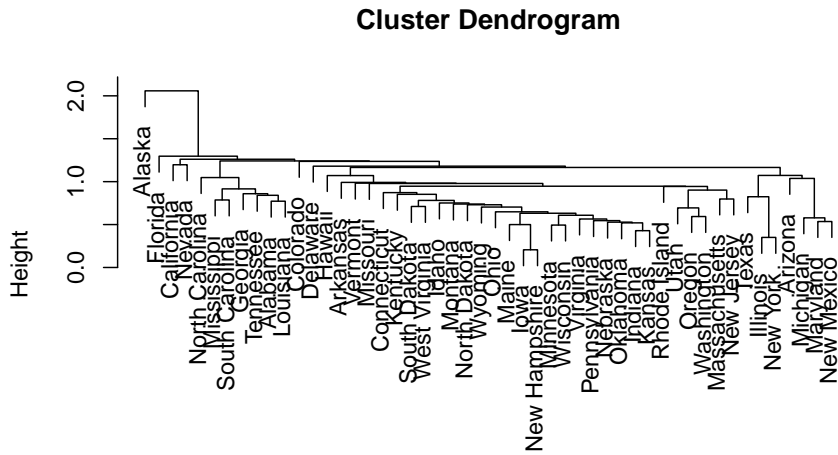
Here the built-in R function `hclust()` is used:

3.4.1 Compute pairwise distance matrices

```
dist.out <- dist(df, method = "euclidean")
```

3.4.2 Single Linkage

```
out.single.euc <- hclust(daisy(df, metric="euclidean"), method="single")
# try other metric="euclidean"
plot(out.single.euc)
```

```
daisy(df, metric = "euclidean")
hclust(*, "single")
```

```
# decide to cut the tree at height 1
out.single.euc <- cutree(out.single.euc, h=1.5)
# view cluster allocation
names(out.single.euc) <- rownames(df)
sort(out.single.euc)
```

##	Alabama	Arizona	Arkansas	California	Colorado
##	1	1	1	1	1
##	Connecticut	Delaware	Florida	Georgia	Hawaii
##	1	1	1	1	1
##	Idaho	Illinois	Indiana	Iowa	Kansas
##	1	1	1	1	1
##	Kentucky	Louisiana	Maine	Maryland	Massachusetts
##	1	1	1	1	1
##	Michigan	Minnesota	Mississippi	Missouri	Montana
##	1	1	1	1	1
##	Nebraska	Nevada	New Hampshire	New Jersey	New Mexico
##	1	1	1	1	1
##	New York	North Carolina	North Dakota	Ohio	Oklahoma
##	1	1	1	1	1
##	Oregon	Pennsylvania	Rhode Island	South Carolina	South Dakota
##	1	1	1	1	1
##	Tennessee	Texas	Utah	Vermont	Virginia
##	1	1	1	1	1


##	Washington	West Virginia	Wisconsin	Wyoming	Alaska
##	1	1	1	1	2

3.4.3 Complete Linkage

```
hc <- hclust(dist.out, method = "complete")
```

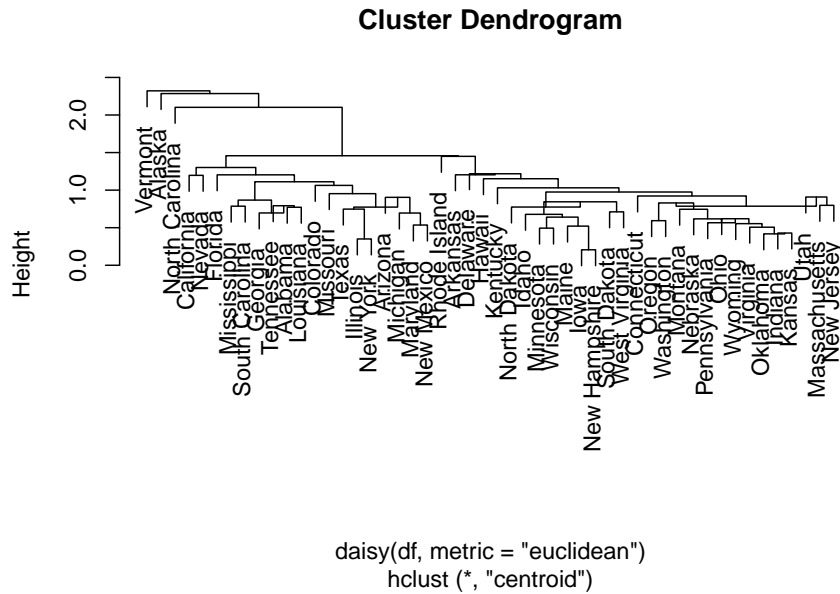
Visualization of hclust

```
plot(hc, labels = F, -1)
rect.hclust(hc, k = 2, border = 2:3) # Add rectangle around 2 clusters, try with 3?
```

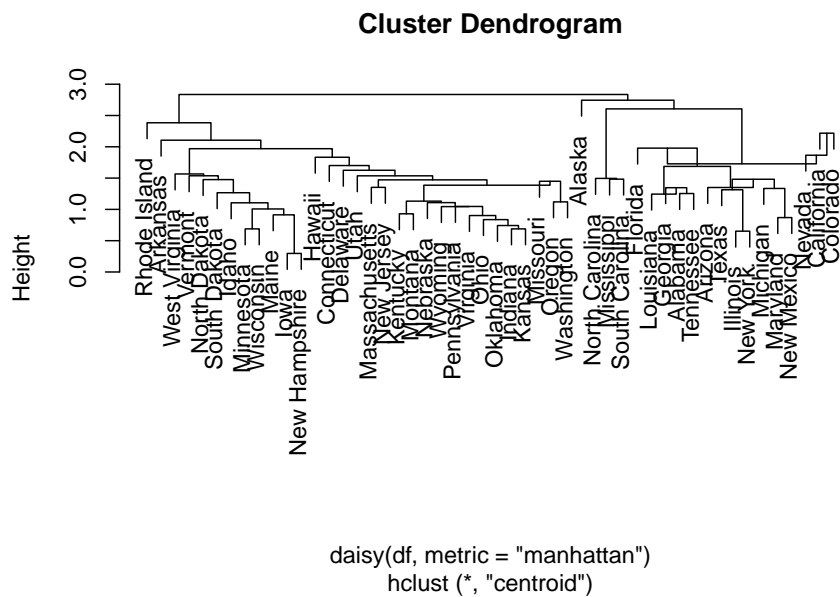


3.4.4 Centroid

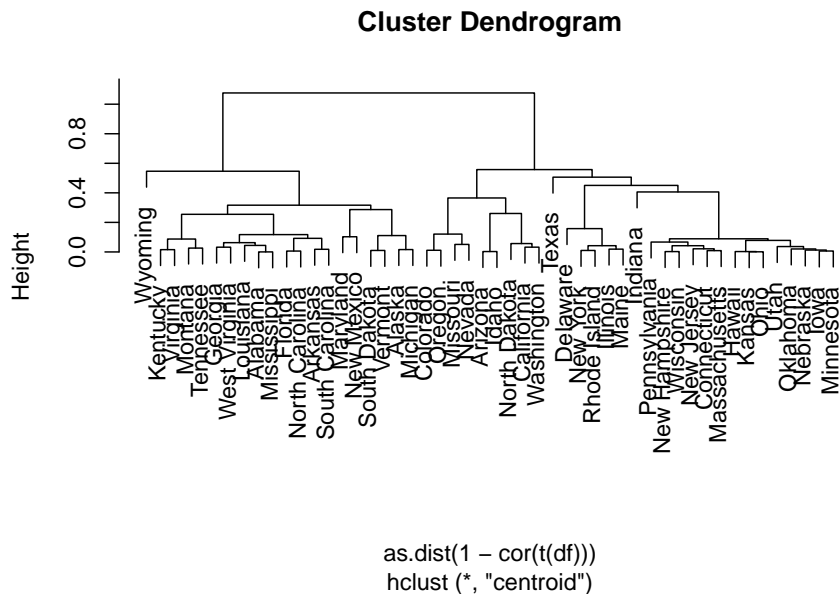
```
# Centroid clustering
out.centroid.euc <- hclust(daisy(df, metric="euclidean"), method="centroid")
plot(out.centroid.euc)
```



```
out.centroid.city <- hclust(daisy(df,metric="manhattan"),method="centroid")
plot(out.centroid.city)
```



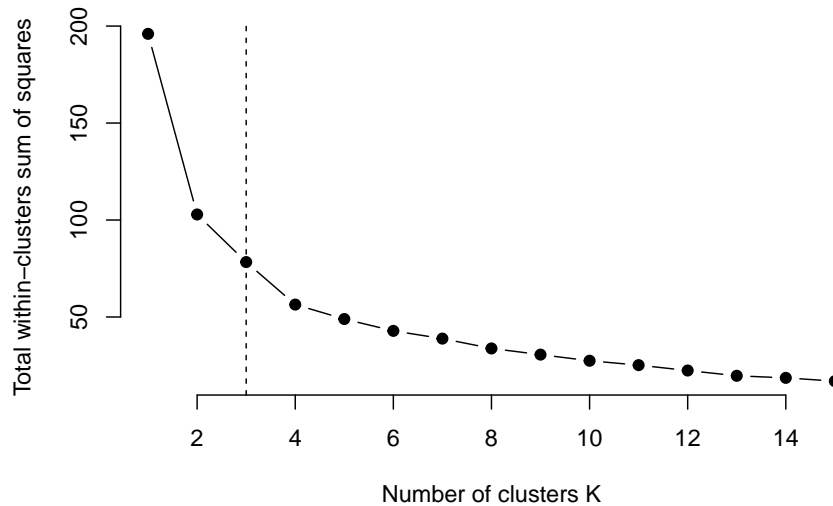
```
out.centroid.cor <- hclust(as.dist(1-cor(t(df))),method="centroid")
plot(out.centroid.cor)
```



3.5 Methods for determining number of clusters

3.5.1 Elbow method for k-means clustering

```
set.seed(123)
# Compute and plot wss for k = 2 to k = 15
k.max <- 15 # Maximal number of clusters
df.out <- df
wss <- sapply(1:k.max,
function(k){kmeans(df.out, k, nstart=10 )$tot.withinss})
plot(1:k.max, wss, type="b", pch = 19, frame = FALSE, xlab="Number of clusters K", ylab="Within-cluster sum of squares")
abline(v = 3, lty =2)
```

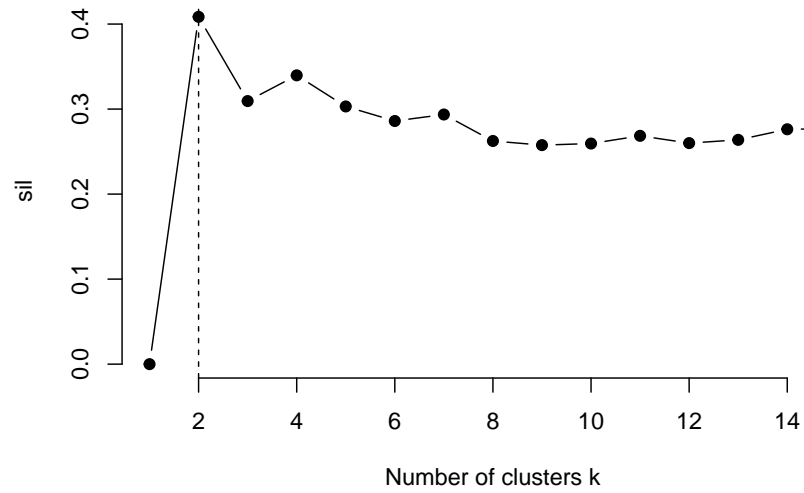


According to the elbow method, the optimal number of clusters suggested for the K-means algorithm is 3.

3.5.2 Average silhouette method for k-means clustering

```
k.max <- 15
data.out <- df
sil <- rep(0, k.max)
# Compute the average silhouette width for
# k = 2 to k = 15
for(i in 2:k.max){
  km.res <- kmeans(df.out, centers = i, nstart = 25)
  ss <- silhouette(km.res$cluster, dist(df.out))
  sil[i] <- mean(ss[, 3])
}
```

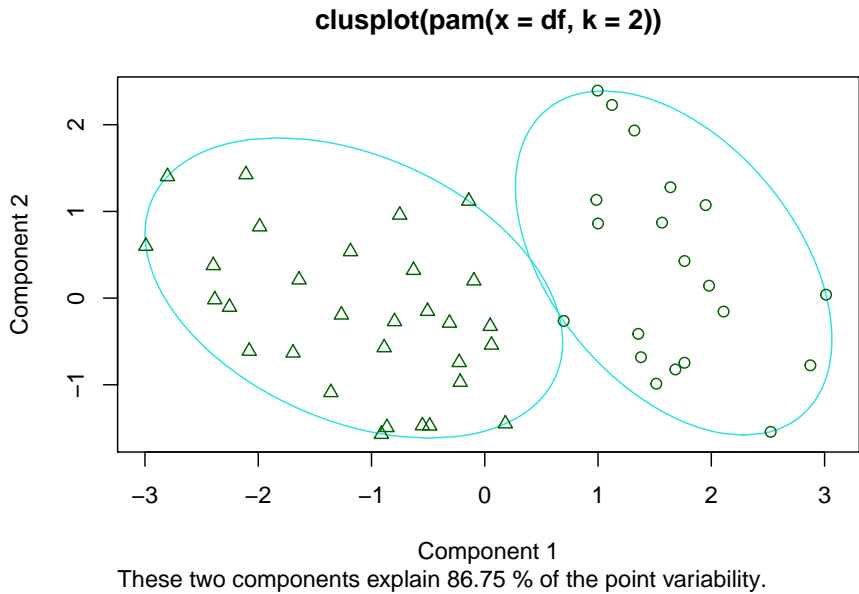
```
# Plot the average silhouette width
plot(1:k.max, sil, type = "b", pch = 19,
     frame = FALSE, xlab = "Number of clusters k")
abline(v = which.max(sil), lty = 2)
```




According to the silhouette method the optimal number of clusters suggested for the Kmeans algorithm is 2.

3.5.3 Average silhouette method for PAM clustering

```
#clusplot(pam.out, main = "Cluster plot, k = 2", color = TRUE)
plot(pam.out)
```



Silhouette plot of pam(x = df, k = 2)



These two components explain 86.75% of the point variability.


This table shows how to use the average silhouette width value:

Range of SC : Interpretation 0.71-1.0 : A strong structure has been found 0.51-0.70 : A reasonable structure has been found 0.26-0.50 : The structure is weak

and could be artificial < 0.25 : No substantial structure has been found
According to the table, the fit is weak.

3.5.4 Average silhouette method for hierarchical clustering

```
plot(silhouette(cutree(hc,2),dist.out))
```



Average silhouette width : 0.4

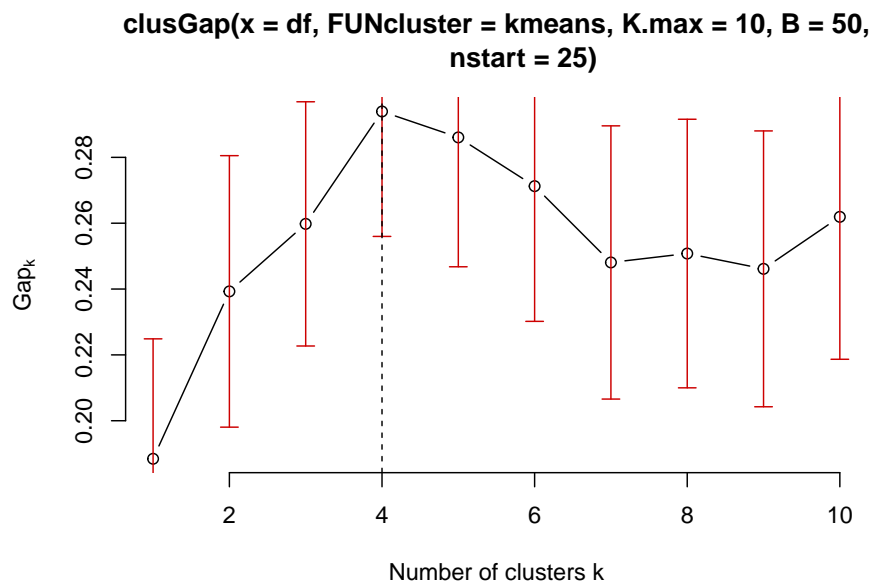
This table shows how to use the average silhouette width value:

Range of SC: Interpretation 0.71-1.0 : A strong structure has been found 0.51-0.70 : A reasonable structure has been found 0.26-0.50 : The structure is weak and could be artificial < 0.25 : No substantial structure has been found

The result for hierarchical clustering is similar to that of PAM. The conclusion we can make is that fit is weak.

3.5.5 Gap Statistic for K means clustering

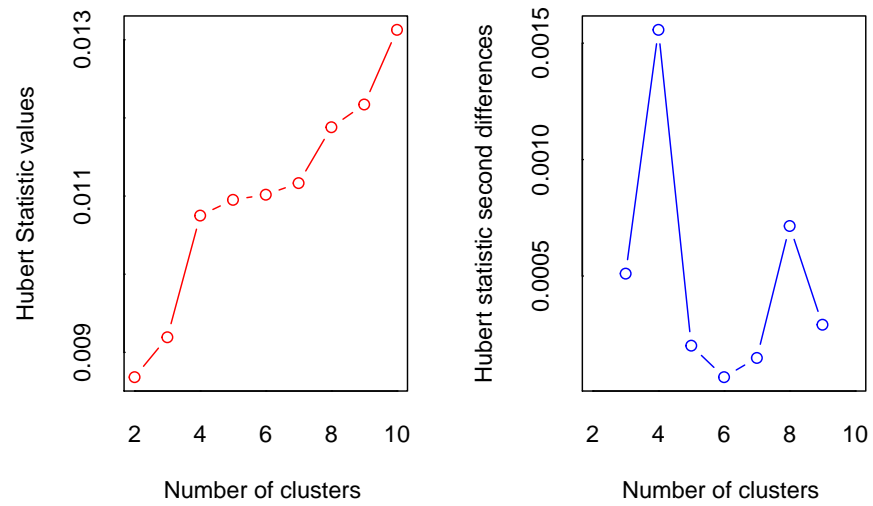

```
# Compute gap statistic
gap_stat <- clusGap(df, FUN = kmeans, nstart = 25, K.max = 10, B = 50)
# Print the result
plot(gap_stat, frame = FALSE, xlab = "Number of clusters k")
abline(v = 4, lty = 2)
```



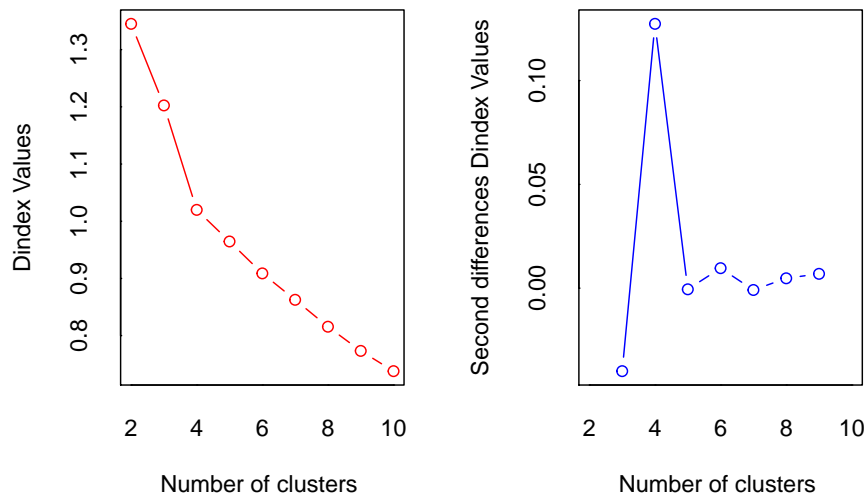
According to the Gap Statistic the 'optimal number of clusters chosen for the Kmeans algorithm is 4!

Using the NbClust package which uses a vote to chose the number of clusters. The following example determine the number of clusters using all statistics:

```
res.nb <- NbClust(df, distance = "euclidean", min.nc = 2, max.nc = 10, method = "complete", index = "all")
```



```
## *** : The Hubert index is a graphical method of determining the number of clusters.
##           In the plot of Hubert index, we seek a significant knee that corresponds to a
##           significant increase of the value of the measure i.e the significant increase in the
##           index second differences plot.
##
```



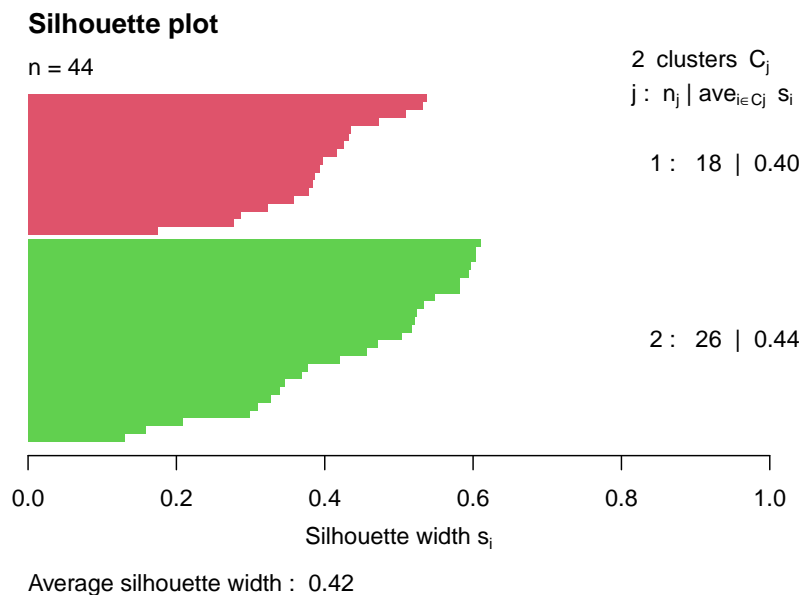
```
## *** : The D index is a graphical method of determining the number of clusters.
##       In the plot of D index, we seek a significant knee (the significant peak in Di
##       second differences plot) that corresponds to a significant increase of the val
##       the measure.
##
## *****
## * Among all indices:
## * 9 proposed 2 as the best number of clusters
## * 4 proposed 3 as the best number of clusters
## * 6 proposed 4 as the best number of clusters
## * 2 proposed 5 as the best number of clusters
## * 1 proposed 8 as the best number of clusters
## * 1 proposed 10 as the best number of clusters
##
##       ***** Conclusion *****
##
## * According to the majority rule, the best number of clusters is 2
##
## *****
```

When all statistics in the NbClust package are allowed to vote, the majority (in this case 9 out of 23) propose that the ‘optimal’ number of clusters should be 2.

3.6 Clustering with CLARA

R function for computing CLARA is found in the `cluster` package.

```
clarax <- clara(df, 2, samples=10)
# Silhouette plot
plot(silhouette(clarax), col = 2:3, main = "Silhouette plot")
```



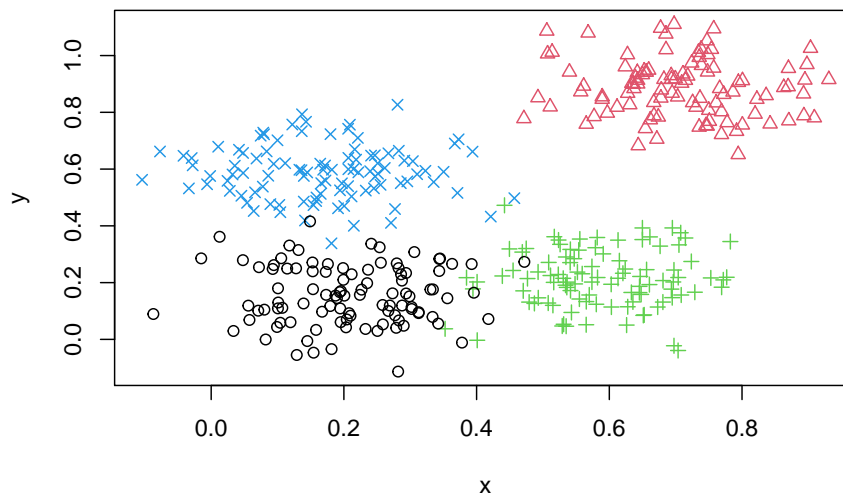
The overall average silhouette width is 0.42 meaning that the fit is weak (see table above showing range for S_i and corresponding interpretation).

3.7 Clustering with DBSCAN

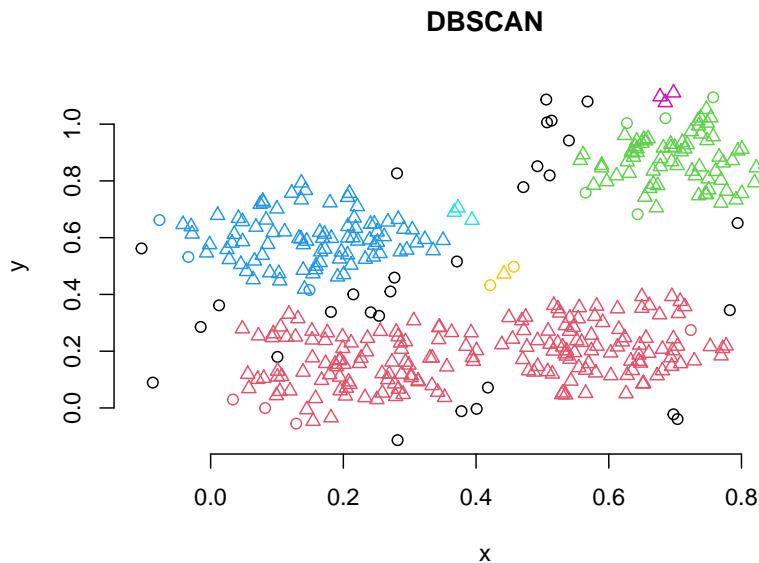
To illustrate the application of DBSCAN we will use a very simple artificial data set of four slightly overlapping Gaussians in two-dimensional space with 100 points each. We set the random number generator to make the results reproducible and create the data set as shown below. The function `dbscan()` is found in the `fpc` package.

```
set.seed(2)
n <- 400
x <- cbind(
  x = runif(4, 0, 1) + rnorm(n, sd = 0.1),
```

```
y = runif(4, 0, 1) + rnorm(n, sd = 0.1)
)
true_clusters <- rep(1:4, time = 100)
plot(x, col = true_clusters, pch = true_clusters)
```



```
# To apply DBSCAN, we need to decide on the neighborhood radius eps
# and the density threshold minPts.
# The rule of thumb for minPts is to use at least the number of
# dimensions of the data set plus one. In our case, this is 3.
db <- fpc::dbscan(x, eps = .05, MinPts = 3 )
# Plot DBSCAN results
plot(db, x, main = "DBSCAN", frame = FALSE)
```

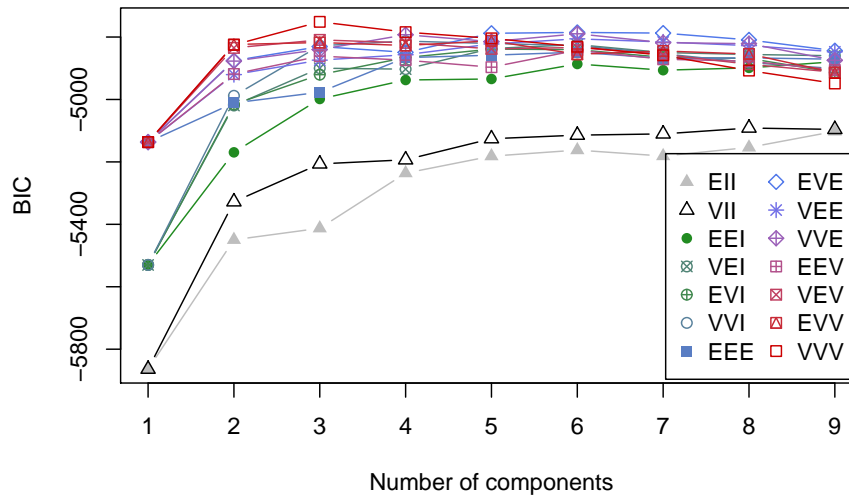


DBSCAN has found three clusters in the data.

3.8 Clustering using mixture models

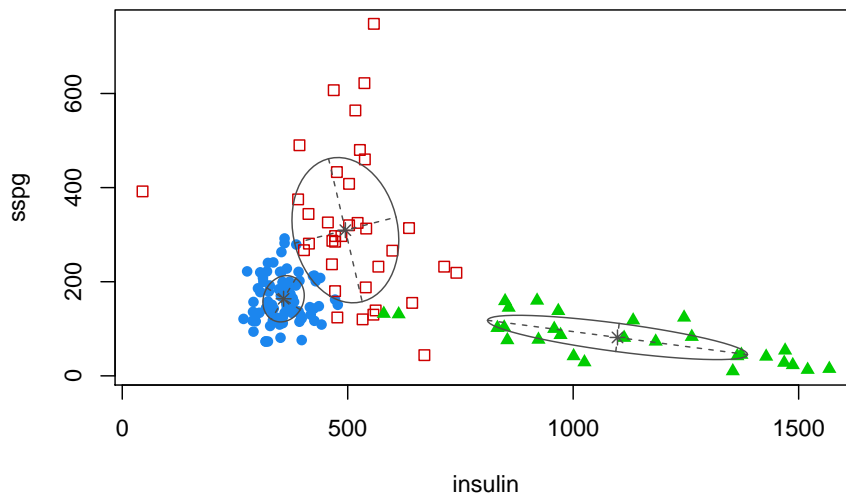
For this you need the function `Mclust()` in the `mclust` package. There are 14 model options available in the R package `mclust`. In one dimension though these collapse into only two models: E for equal variance and V for varying variance. In more dimensions, the model identifiers encode geometric characteristics of the model. For example, EVI denotes a model in which the volume of all clusters are equal (E), the shapes of the clusters may vary (V), and the orientation is the identity (I). That is, clusters in this model have diagonal covariances with orientation parallel to the coordinate axes.

```
library(mclust)
data("diabetes")
# Run the function to see how many clusters
# it finds to be optimal, set it to search for
# at least 1 model and up 20.
d_clust <- Mclust(diabetes[, -1])
plot(d_clust, diabetes[, -1], what="BIC")
```



The plot shows the results of `mclust` for the 10 available model parameterizations and up to 9 clusters for the diabetes dataset. The best model is considered to be the one with the highest BIC among the fitted models.

```
coordProj(diabetes[,-1],dimens = c(2,3),what="classification",classification
=d_clust$classification,parameters = d_clust$parameters)
```

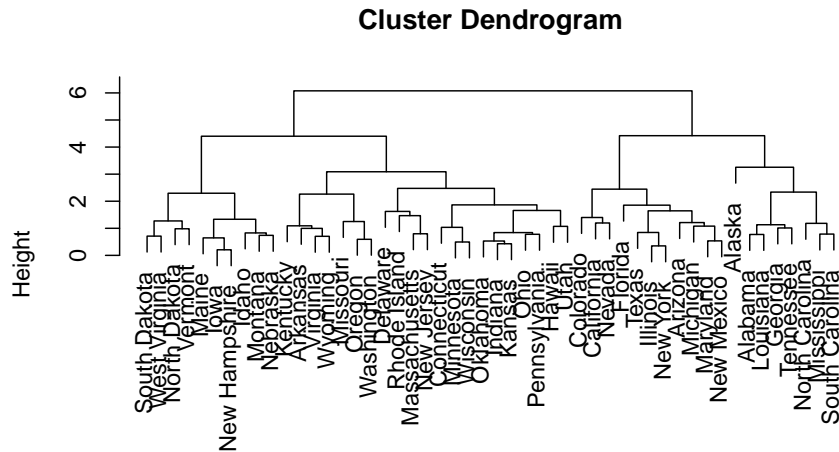


This plot shows the projection of the diabetes data with different symbols indicating the classification corresponding to the best model as determined by `mclust`. The component means are marked and ellipses with axes are drawn corresponding to their covariances. In this case there are three components, each with a different covariance. For more detailed interpretation see (C.Fraley and A.E. Raftery, *Model based Methods of Classification*:

Using the `mclust` Software in Chemometrics. *Journal of Statistical Software*, Vol. 18, 2007)

3.9 Cluster Profiling

```
out.complete.euc <- hclust(daisy(df,metric="euclidean"),method="complete")
plot(out.complete.euc)
```

```
daisy(df, metric = "euclidean")
hclust (*, "complete")
```

```
out.complete.euc <- cutree(out.complete.euc, h=3.1)
clusvec <- out.complete.euc
```

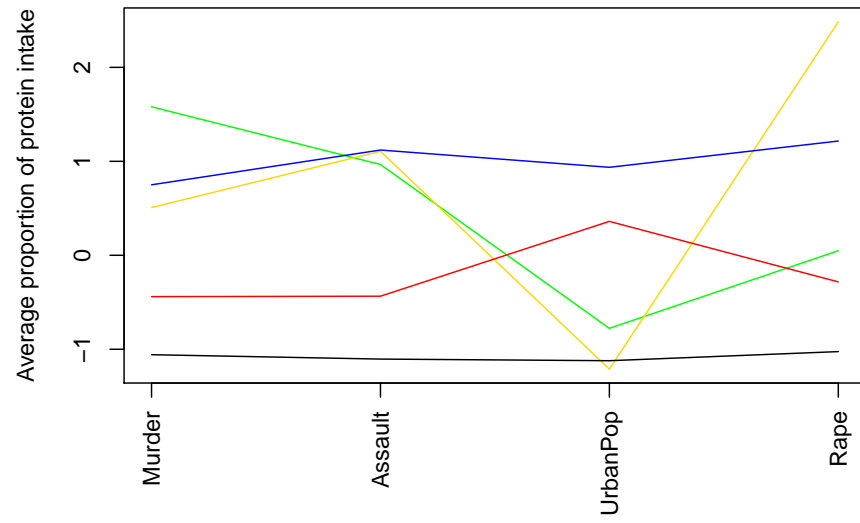
calculate means

```
class.means <- apply(df, 2, function(x) tapply (x, clusvec, mean))
class.means
```

```
##      Murder      Assault      UrbanPop      Rape
## 1  1.5803956  0.9662584 -0.7775109  0.04844071
## 2  0.5078625  1.1068225 -1.2117642  2.48420294
## 3  0.7499801  1.1199128  0.9361748  1.21564322
## 4 -0.4400338 -0.4353831  0.3607592 -0.28303852
## 5 -1.0579703 -1.1046626 -1.1219527 -1.02515543
```

plot means

```
plot (c(1,ncol(df)),range(class.means),type="n",xlab="",ylab="Average proportion of protein intake",
axis (side=1, 1:ncol(df), colnames(df), las=2)
#ensure you list enough colours for the number of clusters
colvec <- c("green","gold","blue","red","black")
for (i in 1:nrow(class.means))
  lines (1:ncol(df),class.means[i,],col=colvec[i])
```



Chapter 4

Self Organising Maps

For this section, please update your browser on the 20th of October Tuesday.